

# TDDC17: Introduction to Automated Planning

Jonas Kvarnström

Automated Planning Group

Department of Computer and Information Science

Linköping University

[jonas.kvarnstrom@liu.se](mailto:jonas.kvarnstrom@liu.se) – 2016

## Introduction to Planning

# One way of defining planning:

*Using **knowledge** about the world,  
including possible actions and their results,  
to **decide** what to do and when  
in order to achieve an **objective**,  
**before** you actually start doing it*

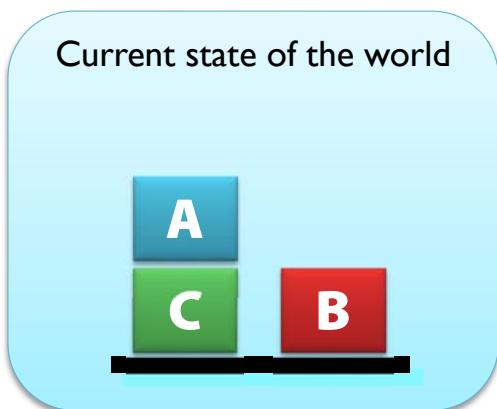
## Domains 1: Blocks World



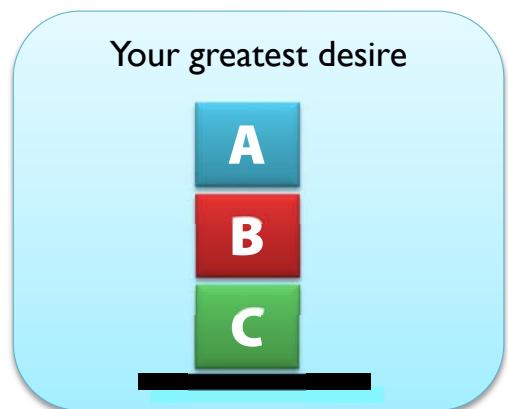
- Classical example: The **Blocks World**



You



Current state of the world

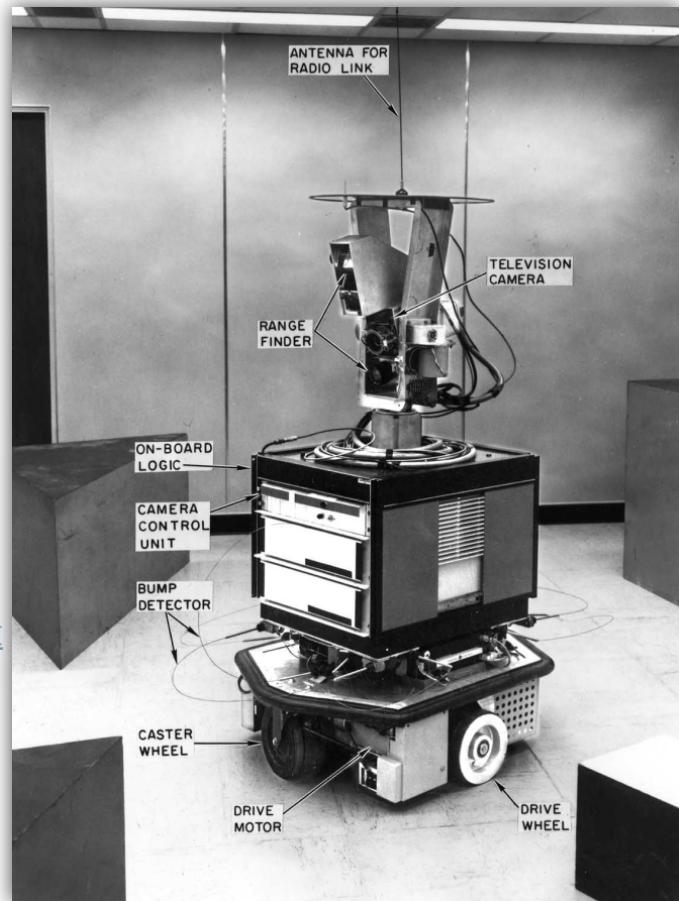


Your greatest desire

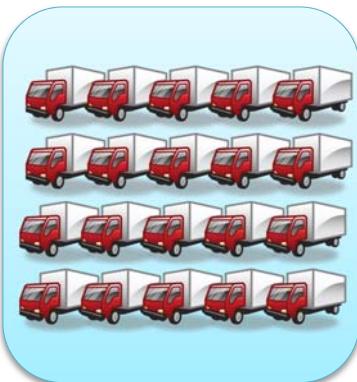
# Domains 2: Shakey

5  
junk@idai

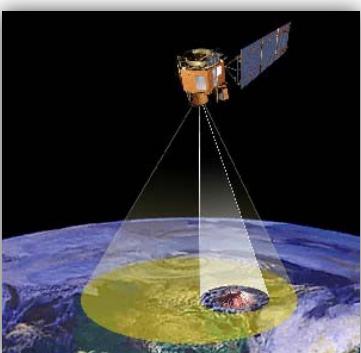
- Classical robot example:  
**Shakey (1969)**
  - Available **actions**:
    - Moving to another location
    - Turning light switches on and off
    - Opening and closing doors
    - Pushing movable objects around
    - ...
  - **Goals**:
    - Be in room 4 with objects A,B,C
    - <http://www.youtube.com/watch?v=qXdn6ynwpiI>



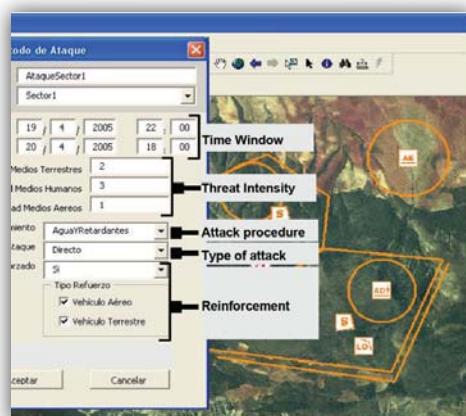
# Domains 3

6  
junk@idai

**Logistics:**  
Use a fleet of trucks  
to efficiently deliver  
packages



**On-board planning**  
to view interesting  
natural events:  
<http://ase.jpl.nasa.gov/>



**SIADEX –**  
**plan for firefighting**  
Limited resources  
Plan execution is  
dangerous!

# Domains 4: Dock Worker Robots (DWR)



**Problem description**



**Solver:**  
Planning Algorithm



**Solution**

Could be a **customized solver**

```

plan      = [ ];
s        = current state of the world;
while (exists b1,b2 [ s.isOn(b1,b2) ]):
    plan  += "unstack(b1,b2)"
    s     = apply(unstack(b1,b2),s)
    plan  += "putdown(b1)"
    s     = apply(putdown(b1),s)
while (exists b1,b2 [ goal.isOn(b1,b2) & !s.isOn(b1,b2) &
                     s.isClear(b1) & s.isClear(b2) ]):
    plan  += "pickup(b1)"
    s     = apply(pickup(b1),s)
    plan  += "stack(b1,b2)"
    s     = apply(stack(b1,b2),s)
return plan
  
```



**Tear down all towers**

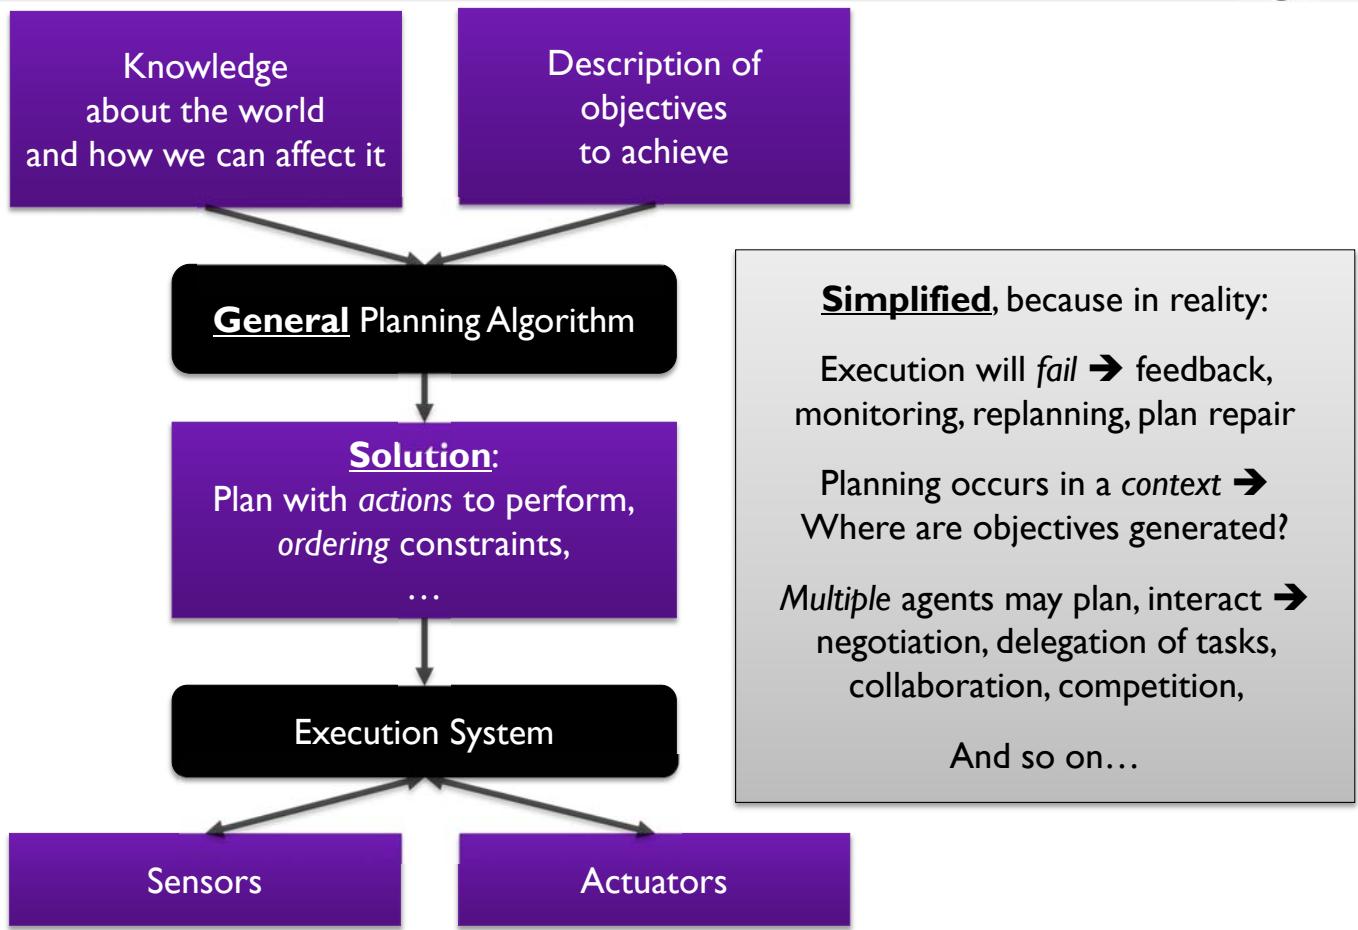
**Rebuild in the right order**

Efficient plans → more complex solvers  
Complex domains → more complex solvers

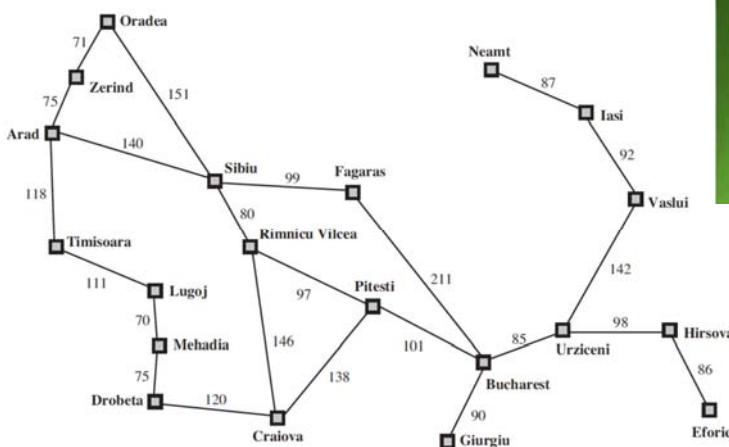
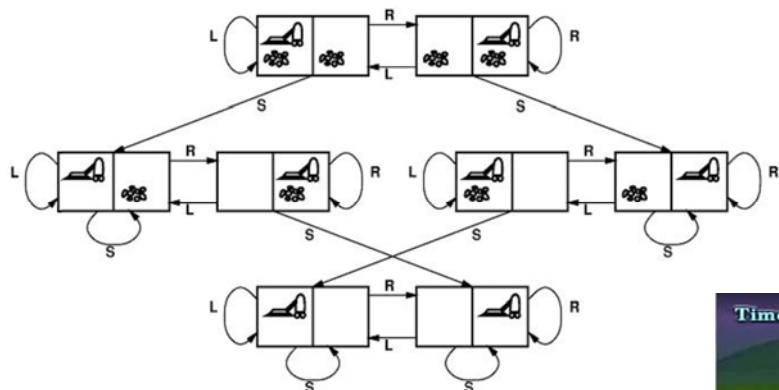
Programming is time-consuming  
Problem changes → more programming required

**We want general + efficient algorithms!**

# AI Planning: A Simplified View



You have already planned – using general search algorithms!

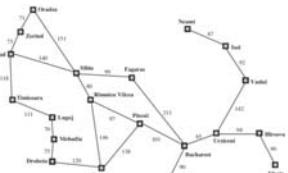


But the representation was problem-specific...

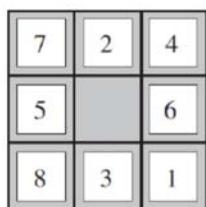


**States:** triple  $(x, y, z)$  with  $0 \leq x, y, z \leq 3$ , where  $x$ ,  $y$ , and  $z$  represent the number of missionaries, cannibals and boats currently on the original bank.

And so was the search guidance!



Straight line distance from city  $n$  to goal city  $n'$



$h_2(n)$ : The sum of the manhattan distances for each tile that is out of place.

$(3+1+2+2+2+3+3+2=18)$ . The manhattan distance is an under-estimate because there are tiles in the way.

How to generalize?

## Classical Planning



First requirement: A formal language for planning problems!

- We want a general language:
  - Allows a wide variety of domains to be modeled
- We want good algorithms
  - Generate plans quickly
  - Generate high quality plans
- Easier with more limited languages



Conflicting desires!

- Many early planners made similar tradeoffs
  - Later called "classical planning"
  - Restricted, but a good place to start

# Modeling Classical Planning Problems

## Planning Domain, Problem Instance



Split knowledge into two parts

### Planning Domain

- General properties of DWR problems
  - There are *containers*, *cranes*, ...
  - Each object has a *location*
  - Possible actions:  
Pick up container, put down container, drive to location, ...

### Problem Instance

- Specific problem to solve
  - *Which* containers and cranes exist?
  - *Where* is everything?
  - *Where should* everything be?  
(More general:  
*What should we achieve?*)

# Objects 1: First-Order → Objects Exist



- Many planning languages have their roots in logic
  - Great – you've already seen this!

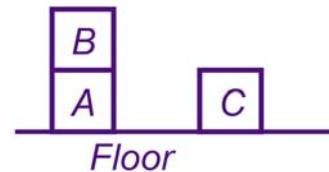
- Propositional syntax

- Plain **propositional facts**:  
~B1,1 , ~S1,1 , OK1,1 , OK2,1 , OK1
  - Used in some planners to simplify parsing, ...

1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2	2,2	3,2	4,2
OK			
1,1	A	2,1	3,1
OK		OK	4,1

- First-order syntax

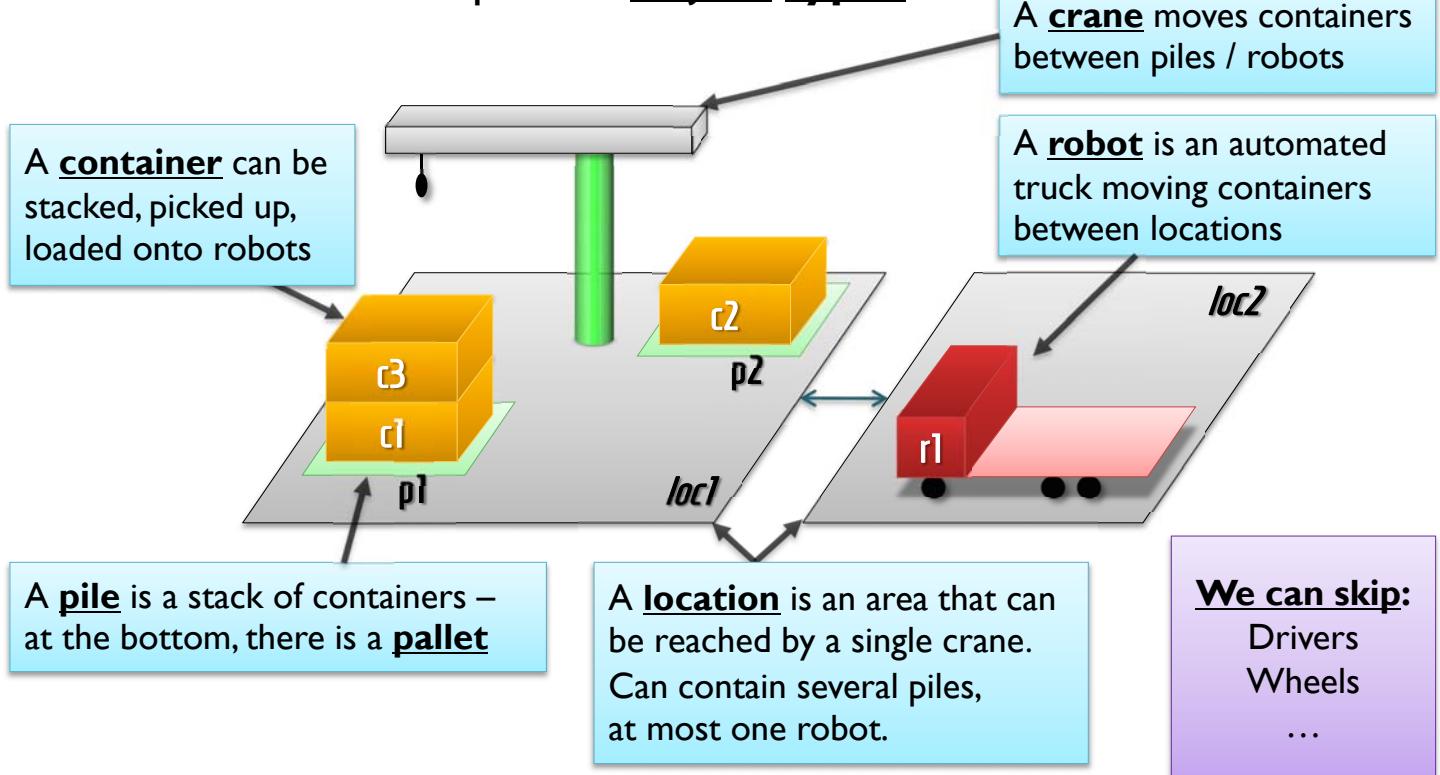
- Objects** and **predicates** (relations):  
On(B,A), On(C,Floor), Clear(B), Clear(C)
  - Used in most implementations



# Objects 2: Object Types



- The domain often specifies object types



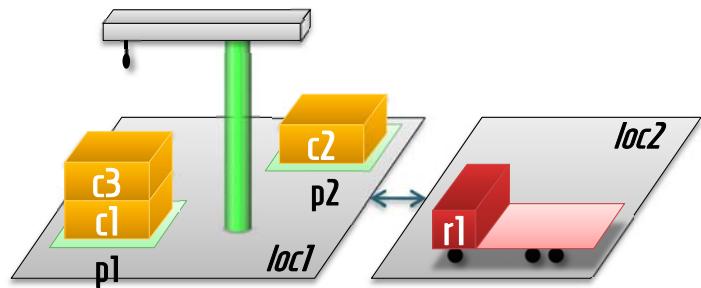
Essential: Determine what is relevant for the problem and objective!

# Objects 3: In the problem instance



## ■ Objects are generally specified in the problem instance

- robot: { r1 }
- location: { loc1, loc2 }
- crane: { k1 }
- pile: { p1, p2 }
- container: { c1, c2, c3, pallet }

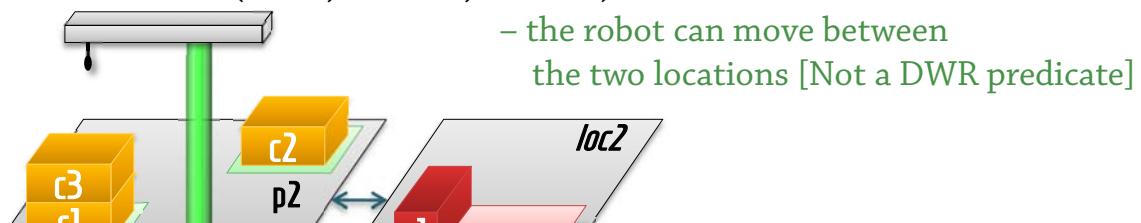


# Facts



## ■ In first-order representations of planning problems:

- Any fact is represented as a logical atom: Predicate symbol + arguments
- Properties of the world
  - **raining** – it is raining [not a DWR predicate!]
- Properties of single objects...
  - **empty(crane)** – the crane is not holding anything
- Relations between objects
  - **attached(pile, location)** – the pile is in the given location
- Relations between >2 objects
  - **can-move(robot, location, location)**
    - the robot can move between the two locations [Not a DWR predicate]



Essential: Determine what is relevant for the problem and objective!

# Facts / Predicates in DWR



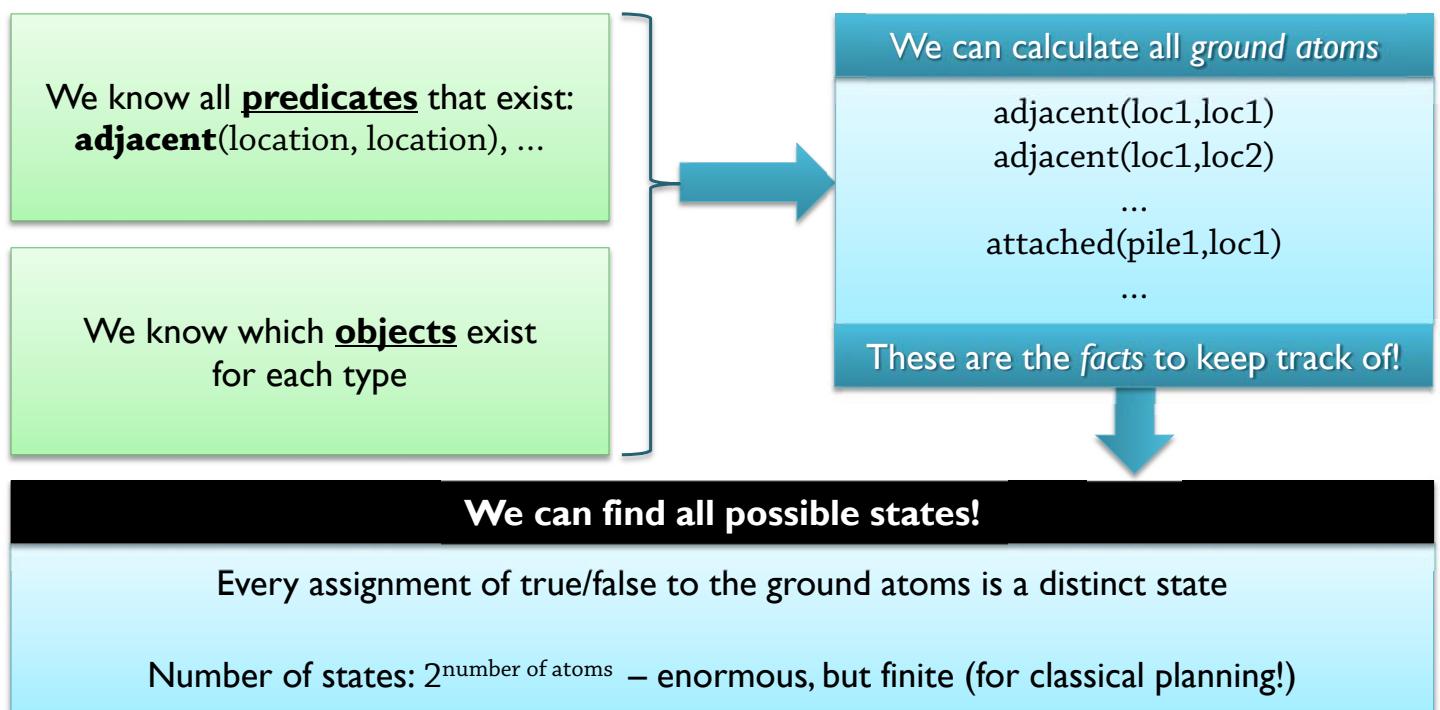
- All predicates for DWR, and their intended meaning:

"Fixed/Rigid" (can't change)	<b>adjacent</b> ( <i>loc1, loc2</i> ) <b>attached</b> ( <i>p, loc</i> ) <b>belong</b> ( <i>k, loc</i> )	; can move from <i>loc1</i> directly to <i>loc2</i> ; pile <i>p</i> attached to <i>loc</i> ; crane <i>k</i> belongs to <i>loc</i>
"Dynamic" (modified by actions)	<b>at</b> ( <i>r, loc</i> ) <b>occupied</b> ( <i>loc</i> ) <b>loaded</b> ( <i>r, c</i> ) <b>unloaded</b> ( <i>r</i> )	; robot <i>r</i> is at <i>loc</i> ; there is a robot at <i>loc</i> ; robot <i>r</i> is loaded with container <i>c</i> ; robot <i>r</i> is empty
	<b>holding</b> ( <i>k, c</i> ) <b>empty</b> ( <i>k</i> )	; crane <i>k</i> is holding container <i>c</i> ; crane <i>k</i> is not holding anything
	<b>in</b> ( <i>c, p</i> ) <b>top</b> ( <i>c, p</i> ) <b>on</b> ( <i>c1, c2</i> )	; container <i>c</i> is somewhere in pile <i>p</i> ; container <i>c</i> is on top of pile <i>p</i> ; container <i>c1</i> is on container <i>c2</i>

## States 1: State of the World



- A **state (of the world)** should specify exactly which facts (**ground atoms**) are true/false in the world at a given time

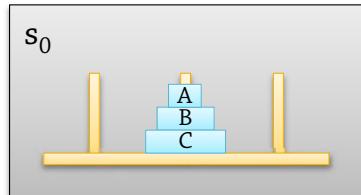


# States 2: Efficient Representation



## ■ Efficient specification and storage for a single state:

- Specify which atoms are true
  - All other atoms have to be false – what else would they be?
- A state of the world is specified as a set containing all variable-free atoms that [are, were, will be] true in the world
  - $s_0 = \{ \text{on}(A,B), \text{on}(B,C), \text{in}(A,2), \text{in}(B,2), \text{in}(C,2), \text{top}(A), \text{bot}(C) \}$



$s_0$	top(A)	in(A,2)
	on(A,B)	in(B,2)
	on(B,C)	in(C,2)
		bot(C)

top(A)  $\in s_0 \rightarrow$  top(A) is true in  $s_0$   
top(B)  $\notin s_0 \rightarrow$  top(B) is false in  $s_0$

# States 3: Initial State



## ■ Initial states in classical STRIPS planning:

- We assume *complete information* about the initial state (before any action)

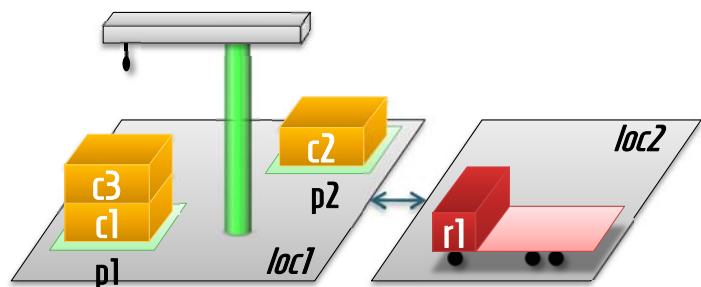
Complete relative to the model:  
We must know everything  
about those predicates and objects  
we have specified...  
But not whether it's raining!

- So we can use a set of true atoms

$\{$

attached(p1, loc1), in(c1, p1), on(c1, pallet), in(c3, p1), on(c3, c1), top(c3, p1),  
attached(p2, loc1), in(c2, p2), on(c2, pallet), top(c2, p2),  
belong(crane1, loc1), empty(crane1),  
at(r1, loc2), unloaded(r1), occupied(loc2),  
adjacent(loc1, loc2), adjacent(loc2, loc1),

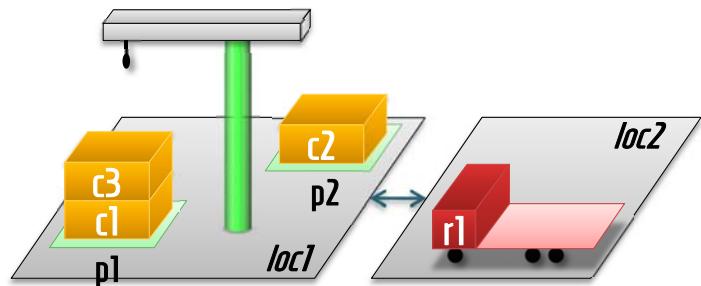
$\}$



# States 4: Goal States



- Classical STRIPS planning: Reach one of possibly many **goal states**
  - Can be specified as a **set of literals** that must hold
  - Example: **Containers 1 and 3 should be in pile 2**
    - We don't care about their order, or any other fact
    - { in(c1,p2), in(c3,p2) }

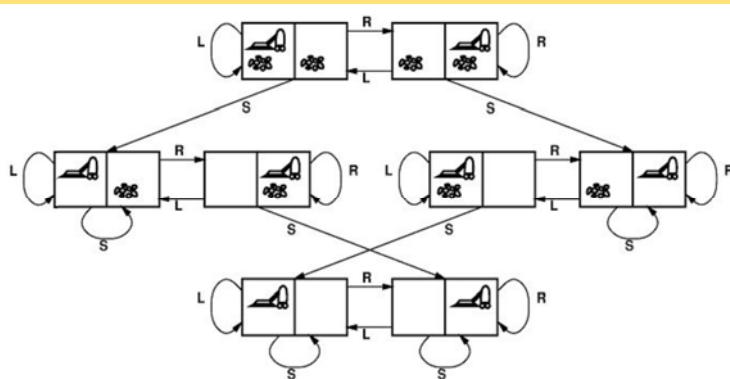


# Actions 1: Intro



- Actions in **plain search** (lectures 2-3):
  - Assumed a *transition / successor function*

**Result(State,Action)** - A description of what each action does (Transition function)



- But how to **specify** it **succinctly**?

# Actions 2: Operators



## ■ Define operators or action schemas:

- **move**(robot, location1, location2)

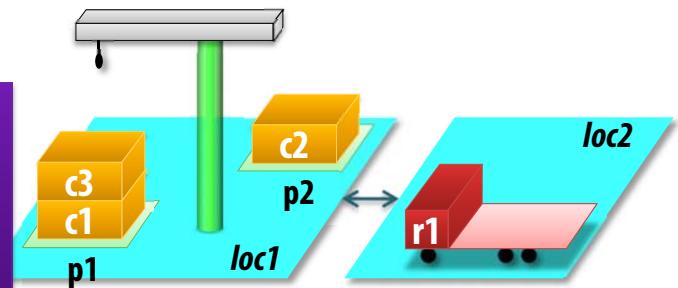
- Precondition: **at**(robot, location1)  $\wedge$   
**adjacent**(location1, location2)  $\wedge$   
 $\neg$  **occupied**(location2)

The action is **applicable**  
in a state  $s$   
if its precond is true in  $s$

- Effects:  
 $\neg$  **at**(robot, location1),  
**at**(robot, location2),  
 $\neg$  **occupied**(location1),  
**occupied**(location2)

The result of applying the  
action in state  $s$ :  
 $s - \text{negated effects}$   
+ positive effects

**Classical planning:**  
Known initial state, known state update function  
→ **deterministic**, can completely predict the  
state of the world after a sequence of actions!



# Actions 3: Instances



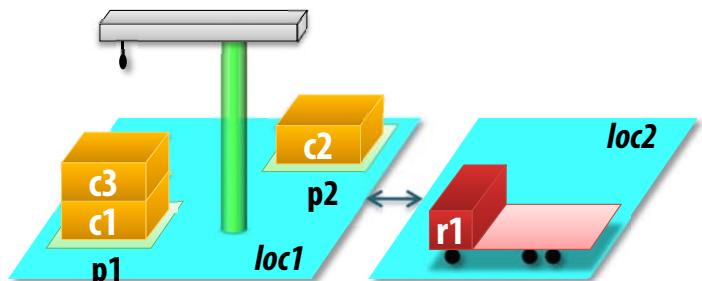
## ■ The planner instantiates these schemas

- Applies them to any combination of parameters  
of the correct type

- **Example: move(r1, loc1, loc2)**

- Precondition: **at**(r1, loc1)  $\wedge$   
**adjacent**(loc1, loc2)  $\wedge$   
 $\neg$  **occupied**(loc2)

- Effects:  
 $\neg$  **at**(r1, loc1),  
**at**(r1, loc2),  
 $\neg$  **occupied**(loc1),  
**occupied**(loc2)



# Actions 4: Step by Step

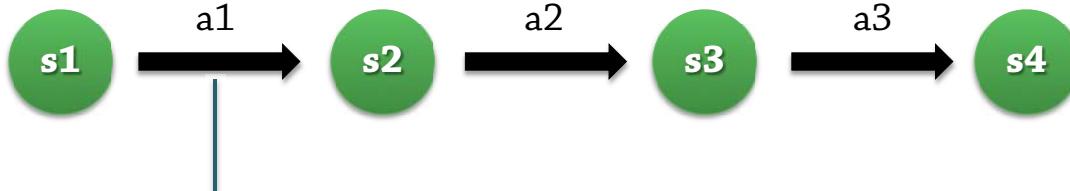


- In classical planning (the basic, limited form):

We know  
the initial  
state

Each action corresponds to one state update

Time is not modeled,  
and never multiple state updates for one action



We know how states are changed by actions

→ Deterministic, can completely predict the state of the world after a sequence of actions!

The solution to the problem  
will be a sequence of actions

## Plan Generation

# With or Without Search



- Important distinction:
  - **Planning** means deciding in advance which actions to perform in order to achieve a goal
  - **Search** will often be a useful tool, but you can get by without it

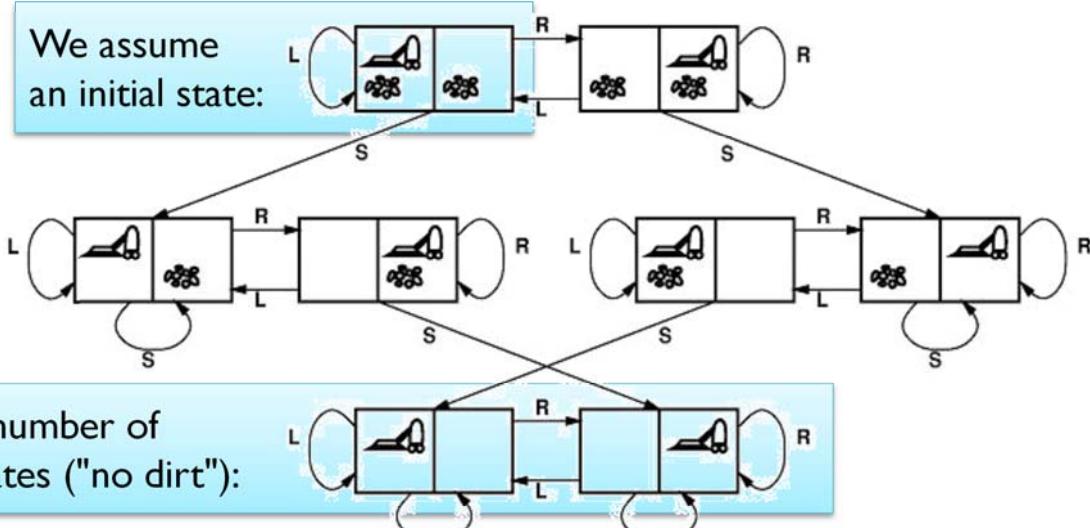
```
plan          = [ ];
s             = current state of the world;
while (exists b1,b2 [ s.isOn(b1,b2) ]):
    plan      += "unstack(b1,b2)"
    s          = apply(unstack(b1,b2), s)
    plan      += "putdown(b1)"
    s          = apply(putdown(b1), s)
while (exists b1,b2 [ goal.isOn(b1,b2) & !s.isOn(b1,b2) &
                     s.isClear(b1) & s.isClear(b2) ]):
    plan      += "pickup(b1)"
    s          = apply(pickup(b1), s)
    plan      += "stack(b1,b2)"
    s          = apply(stack(b1,b2), s)
return plan
```

If we use search:  
What search space?

## Plan Generation Method 1: Forward State Space Search

- Forward state space search: As explored in the Vacuum World

- A search node is simply a world state
- Successor / transition function:
  - One outgoing edge for every executable (applicable) action
  - The action specifies where the edge will lead



Find a path (not necessarily shortest) – **SRS, RSLS, LRLRLSSSRLRS, ...**

How does a state space look?

# State Spaces 1: Towers of Hanoi



Toy Problem I: Towers of Hanoi has a very *regular* state space...



## General Knowledge: Problem Domain

- There are *pegs* and *disks*
- A disk can be *on* a peg
- One disk can be *above* another
- Actions:  
Move topmost disk from *x* to *y*, without placing larger disks on smaller disks

## Specific Knowledge and Objective: Problem Instance

- 3 pegs, 7 disks
- All disks currently on peg 2, in order of increasing size
- We want: All disks on the *third* peg, in order of increasing size

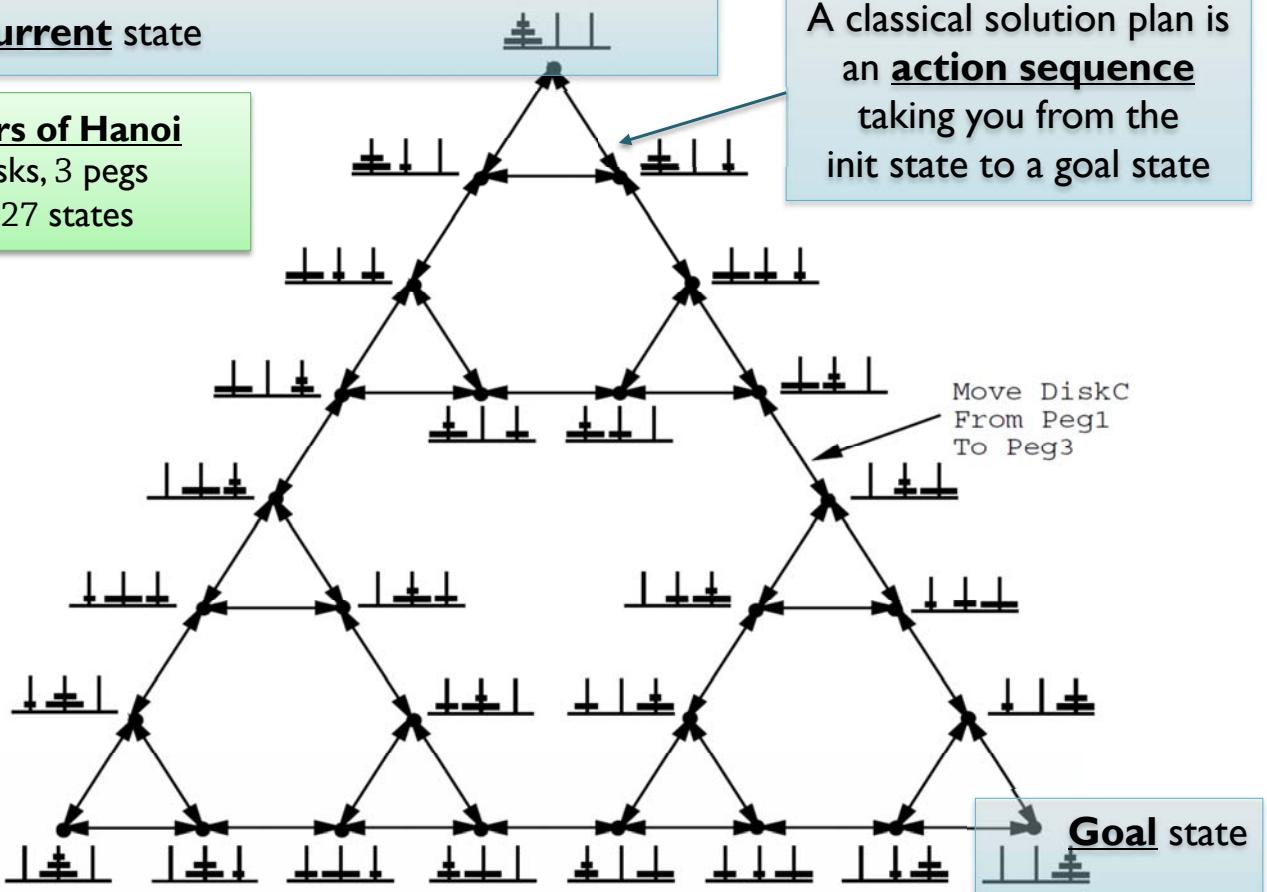
# State Spaces 2: ToH, Actions



## Initial/current state

**Towers of Hanoi**  
3 disks, 3 pegs  
→ 27 states

A classical solution plan is an **action sequence** taking you from the init state to a goal state

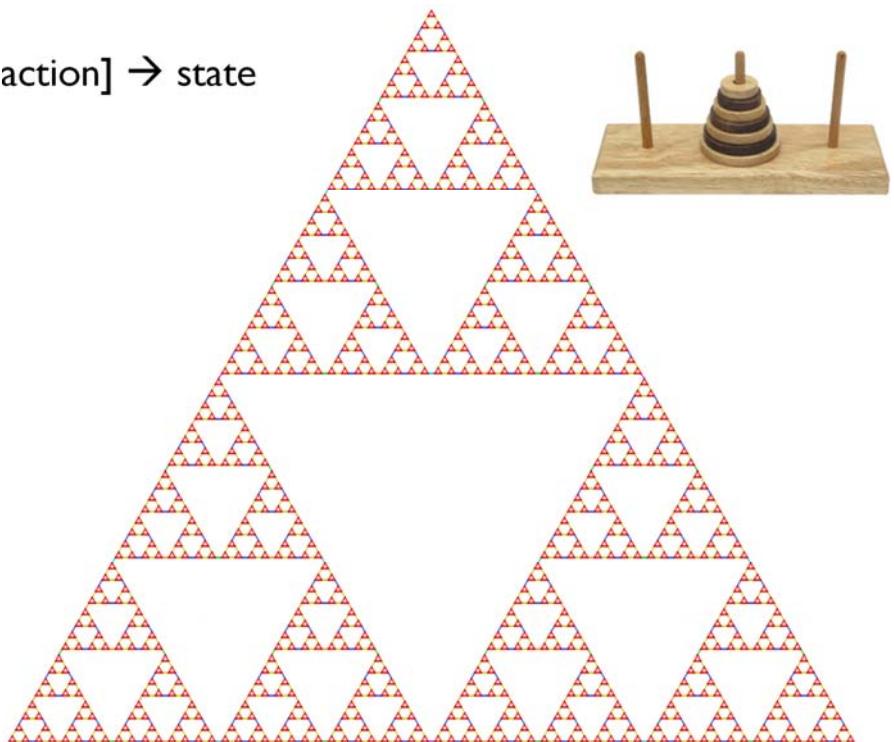


# State Spaces 3: Larger Example



- Larger state space – interesting symmetry

- 7 disks
- 2187 states
- 6558 transitions, [state, action] → state



# State Spaces 4: Blocks World



- A common blocks world version, with 4 operators

- pickup(x) – takes x from the table
- putdown(x) – puts x on the table
- unstack(x, y) – takes x from on top of ?y
- stack(x, y) – puts x on top of y

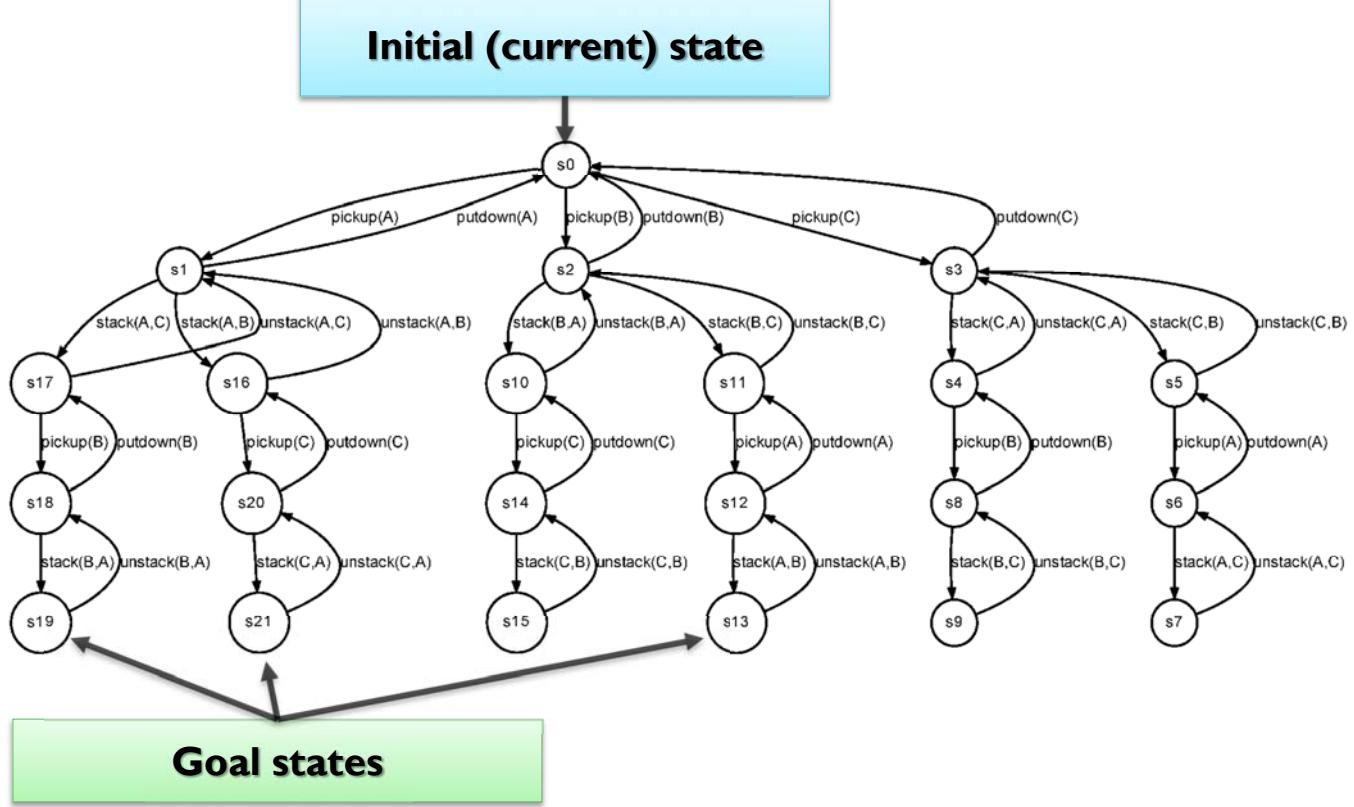


- Predicates (relations) used:

- on(x, y) – block x is on block y
- ontable(x) – x is on the table
- clear(x) – we can place a block on top of x
- holding(x) – the robot is holding block x
- handempty – the robot is not holding any block

clear(A)  
on(A, C)  
ontable(C)  
clear(B) ontable(B)  
clear(D) ontable(D)  
handempty

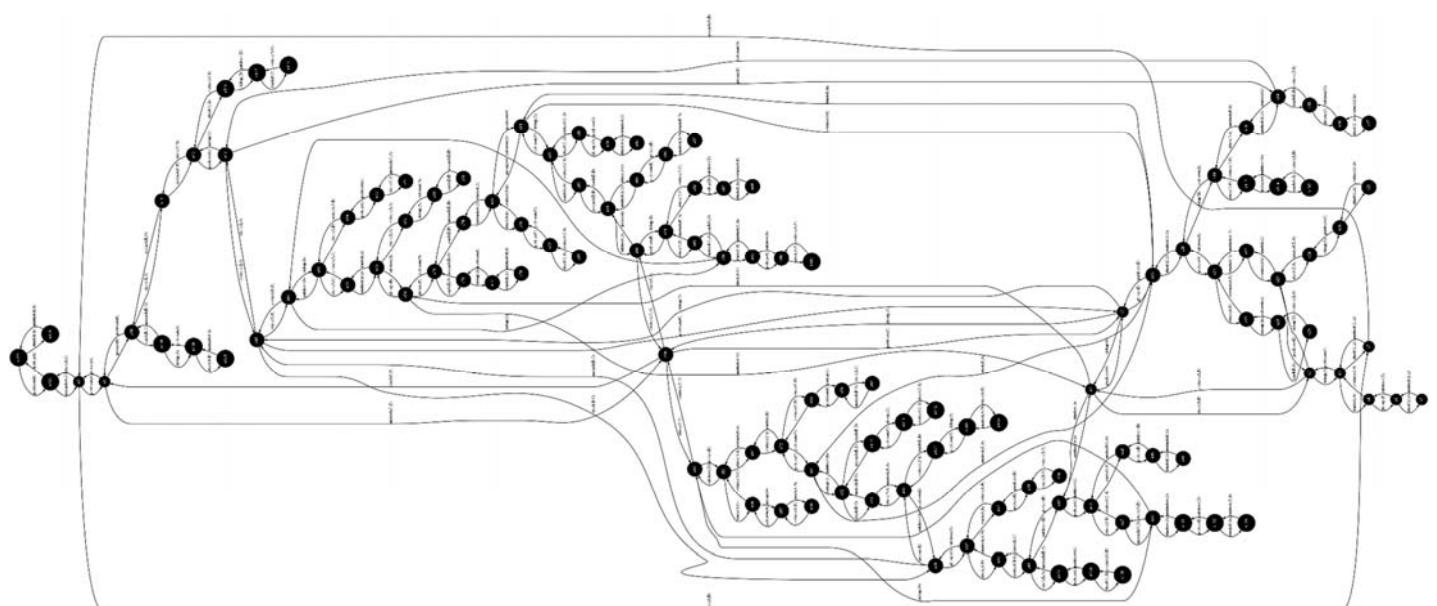
# State Spaces 5: Blocks World, 3 blocks



# State Spaces 6: Blocks World, 4 blocks



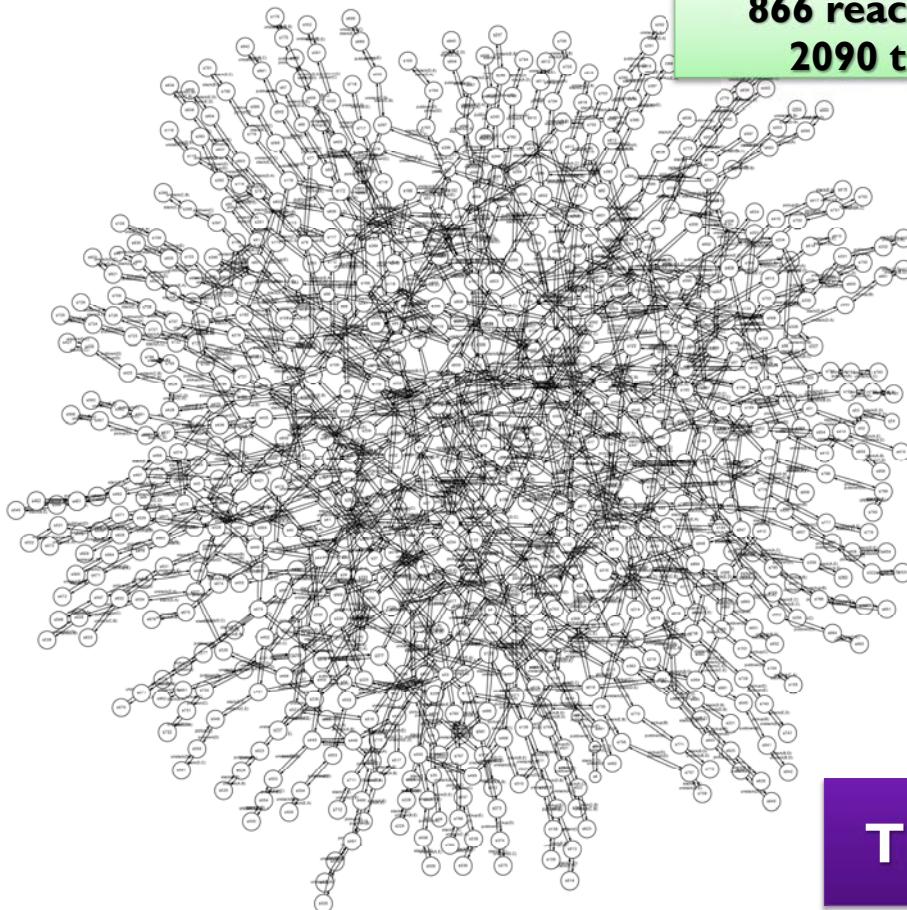
125 reachable states  
272 transitions



# State Spaces 7: Blocks World, 5 blocks



866 reachable states  
2090 transitions



This is tiny!

# State Spaces 8: Reachable State Space

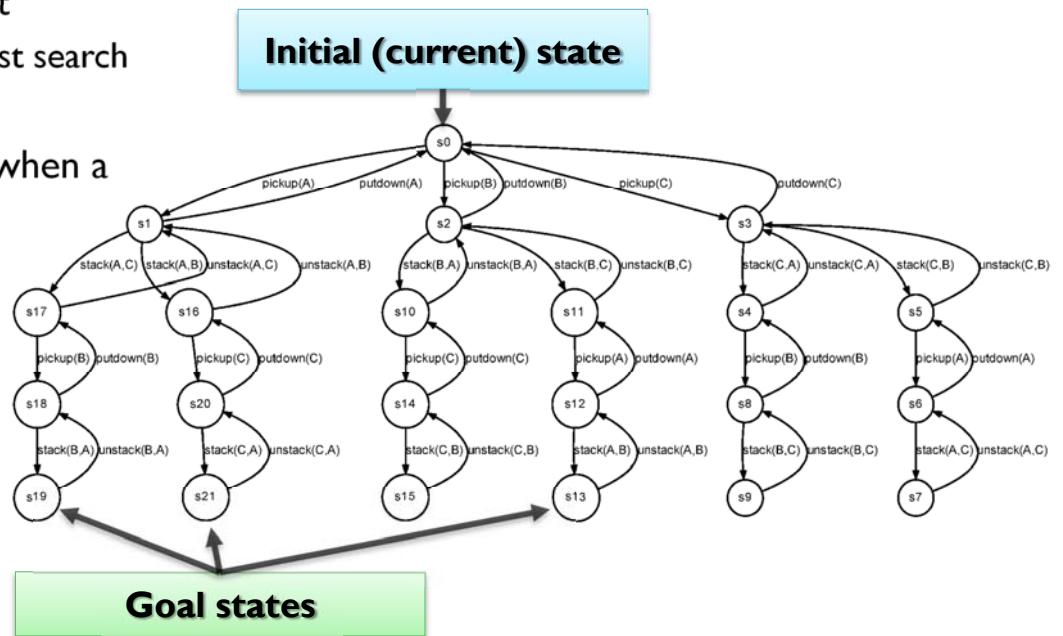


Blocks	States reachable from "all on table"	Transitions (edges) in reachable part
0	1	0
1	2	2
2	5	8
3	22	42
4	125	272
5	866	2090
6	7057	18552
7	65990	186578
8	695417	2094752
9	8145730	25951122
10	...	...
...30	>197987401295571718915006598239796851	

# Forward Search 1



- Forward search:
  - Start in the initial state
  - Apply a **search** algorithm
    - Depth first
    - Breadth first
    - Uniform-cost search
    - ...
  - Terminate when a goal state is found



## Forward Search 2: Step by step



Extremely many states – must generate the graph as we go

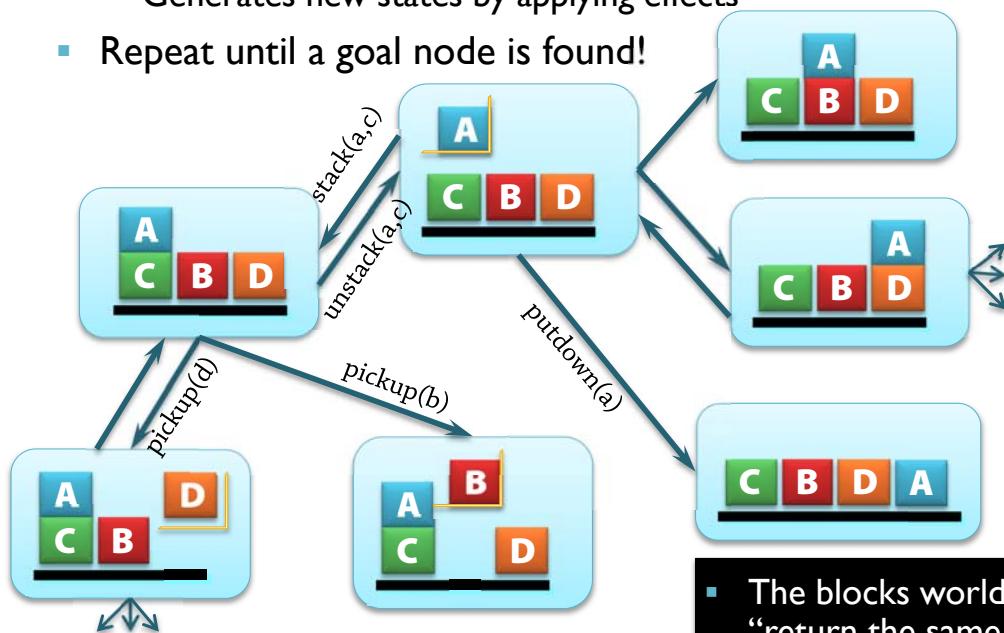
- Generate the initial node == initial state
  - From the initial state *description* in the input



# Forward Search 3: Step by step



- Incremental expansion: Choose a node (using search strategy)
- Expand all possible successors
  - “What actions are applicable in the current state, and where will they take me?”
  - Generates new states by applying effects
- Repeat until a goal node is found!

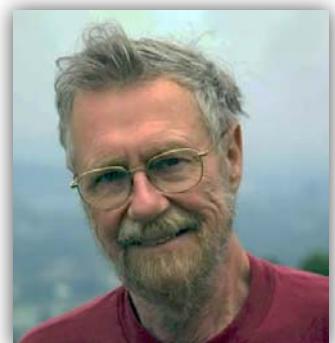


- The blocks world is *symmetric*: Can always “return the same way”
- Not true for all domains!

## Dijkstra



- Which search strategy to use?
  - Breadth first, depth first, iterative deepening depth first, ...
- One possible strategy: Dijkstra's algorithm
  - Uniform-cost search, but terminate when a goal state is found
  - **Efficient**:  $O(|E| + |V| \log |V|)$ 
    - $|V|$  = the number of nodes
    - $|E|$  = the number of edges
  - And generates optimal (cheapest) paths/plans!



Would this algorithm be a solution?

# Dijkstra's Algorithm: Analysis



- Blocks world, 400 blocks initially on the table, goal is a 400-block tower
  - Given that all actions have the same cost, Dijkstra will first consider all plans that stack less than 400 blocks!
    - Stacking 1 block: =  $400 \times 399$  plans, ...
    - Stacking 2 blocks: >  $400 \times 399 \times 398$  plans, ...
  - More than  
163056983907893105864579679373347287756459484163478267225862419762304263994207997664258213955766581163654137118  
16311922048822638316916164832045949028341063579874523269897113293928447980030409667435497403872258873480963719  
24064272436362915472663293976417723601031569414863681933421725283641400148727761800296608761037018087769490614  
84788741874402606226134803936935233568418055950371185351837140548515949431309313875210827888943337113613660928  
31808629961795389295372200673415893327657647047564067391701026030959040303548174221274052329579637773658722452  
549738459404452586503693169340  
0912754853265795909113444084441755664211796  
27432025699299231777374983037  
1882657444844563187930907779661572990289194  
81058521781914647662930023360  
1372350568748665249021991849760646988031691  
39438655119417119333314403154  
1302649432305620215568850657684229678385177  
72535893398611212735245298803  
3087201742432360729162527387508073225578630  
777685901637435541458440833878709344174983977437430327557534417629122448835191721077333875230695681480990867109  
05133210482041360782220646535272711073906611800376194410428900071013695438359094641682253856394743335678545824  
32093210697331749851571100671998530498260475511016725485476618861912891705393547098435020659778689499606904157  
077005797632287669764145095581565056589811721520434612770594950613701730879307727141093526534328671360002096924  
48349430242464906145172664594758586010497684553450747960540890382832026131072217782156434204572434616042404375  
21105232403822580540571315732915984635193126556273109603937188229504400

**1.63 \* 10<sup>1735</sup>**

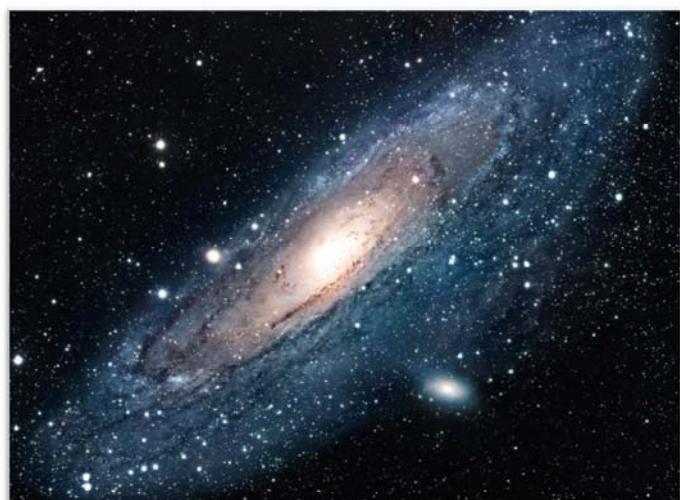
Efficient in terms of the search space size:  $O(|E| + |V| \log |V|)$

The search space is exponential in the size of the input description...

## Fast Computers, Many Cores



- But computers are getting very fast!
  - Suppose we can check  $10^{20}$  states per second
    - >10 billion states per clock cycle for today's computers, each state involving complex operations
  - Then it will only take  $10^{1735} / 10^{20} = 10^{1715}$  seconds...
- But we have multiple cores!
  - The universe has at most  $10^{87}$  particles, including electrons, ...
  - Let's suppose every one is a CPU core
  - → only  $10^{1628}$  seconds  
 $> 10^{1620}$  years
  - The universe is around  $10^{10}$  years old



**Hopeless? No: We need informed search!**

# Heuristic Forward State Space Search

## Domain-Independent Heuristics



### ■ We need:

- A heuristic function  $h(n)$  estimating the cost or difficulty of reaching the goal from node  $n$ 
    - Sometimes, cost = number of actions
    - Sometimes, real costs specified per action
  - A search strategy selecting nodes depending on  $h(n)$
- 
- What's the difference from heuristics in the search lectures?
    - Previously we adapted heuristics to the problem at hand
      - Romania Travel → straight line distance
      - 8-puzzle → pieces out of place, sum of Manhattan distances
    - Now we want to define a general heuristic function
      - Without knowing what planning problem is going to be solved!

# Optimal and Satisficing Planning



## ■ What about admissibility?

- If we must **guarantee** optimal plans:
  - **Need** optimal search strategies such as A\*
  - **Must** use admissible heuristics – or we may get suboptimal plans

## ■ If we are satisfied with a **high-quality plan**:

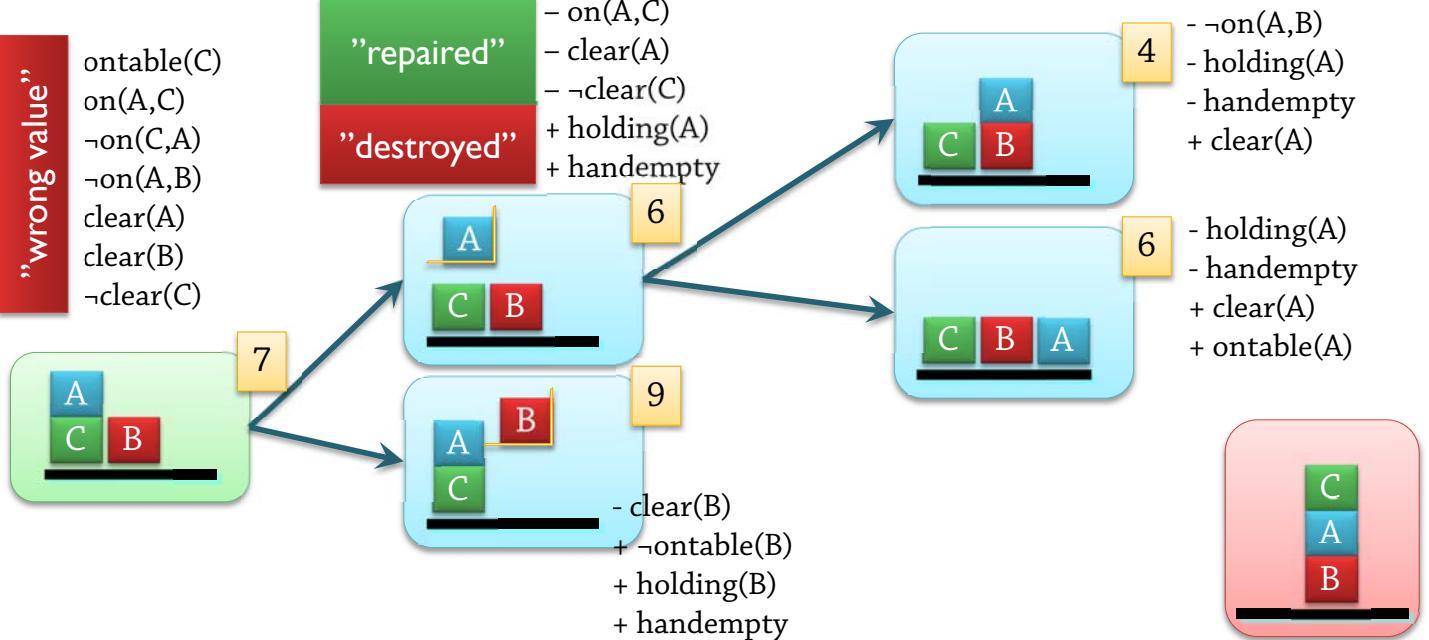
- **Satisficing** planning
  - Find a plan that is sufficiently good, sufficiently *quickly* (where "sufficient" is usually not well-defined)
  - Can use **non-admissible** heuristics, **non-optimal** strategies
    - ...which can often be more **informative** (give us more help in finding reasonably good alternatives)
  - Generally much faster
    - → Handles more complex domains, larger problem instances

# Domain-Independent Heuristics (1)



- A very simple **domain-independent** heuristic:
  - **Count** the number of facts that are “wrong” in the current state  
(similar to “number of pieces out of place”)

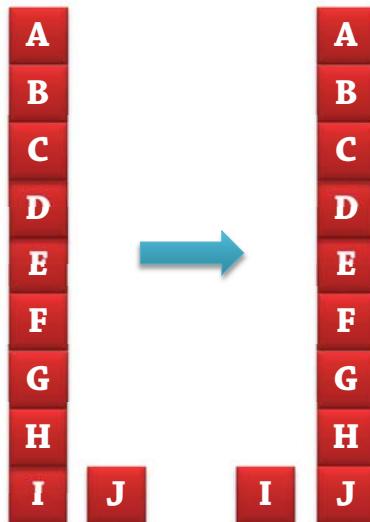
**Optimal:**  
unstack(A,C)  
stack(A,B)  
pickup(C)  
stack(C,A)



# Domain-Independent Heuristics (2)



- Helpful, but not very...



## Optimal solution:

unstack(A,B)	pickup(G)
putdown(B)	stack(G,H)
unstack(B,C)	pickup(F)
putdown(C)	stack(F,G)
unstack(C,D)	pickup(E)
putdown(D)	stack(E,F)
...	...
unstack(H,I)	
stack(H,J)	

**"Wrong facts":**  
**on(H,I), on(H,J),**  
**clear(I), clear(J)**

**Must remove many**  
**"good facts"**  
**to reach the goal**

## Creating Admissible Heuristic Functions: The General Relaxation Principle

How does relaxation apply to planning?

How can we make the definition more precise?

# Fundamental ideas



## We want:

- A way to find an admissible heuristic  $h(s)$  for planning problem  $P$

## One general method:

- Find a problem  $P'$  that we can solve quickly, such that:
  - $\text{cost}(\text{optimal-solution}(P')) \leq \text{cost}(\text{optimal-solution}(P))$

- Solve** problem  $P'$  optimally starting in  $s$ , resulting in solution  $\pi(s)$

- Let  $h(s) = \text{cost}(\pi(s))$

**Or:**  
Find the *cost* of an optimal solution directly from  $P'$ , without actually generating that solution

("sum of Manhattan distances")

## We then know:

- $h(s) = \text{cost}(\pi(s)) = \text{cost}(\text{optimal-solution}(P')) \leq \text{cost}(\text{optimal-solution}(P))$
- $h(s)$  is admissible

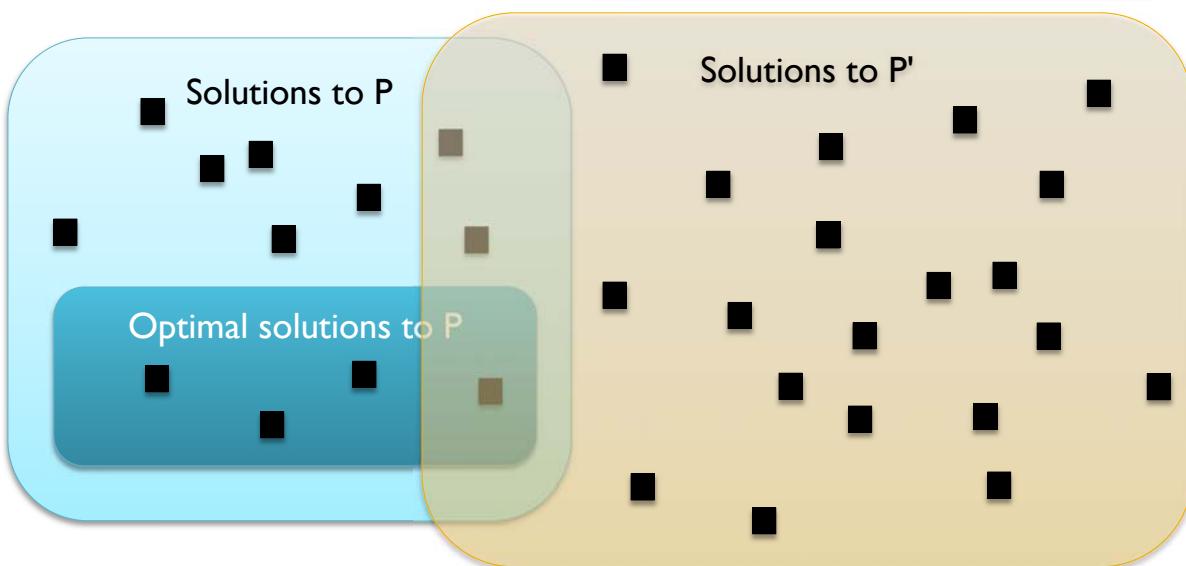
# Fundamental ideas (2)



## How to find a *different* problem with *equal or lower* solution cost?

- Sufficient criterion: One optimal solution to  $P$  remains a solution for  $P'$ 
  - $\text{cost}(\text{optimal-solution}(P')) = \min \{ \text{cost}(\pi) \mid \pi \text{ is any solution to } P' \} \leq \text{cost}(\text{optimal-solution}(P))$

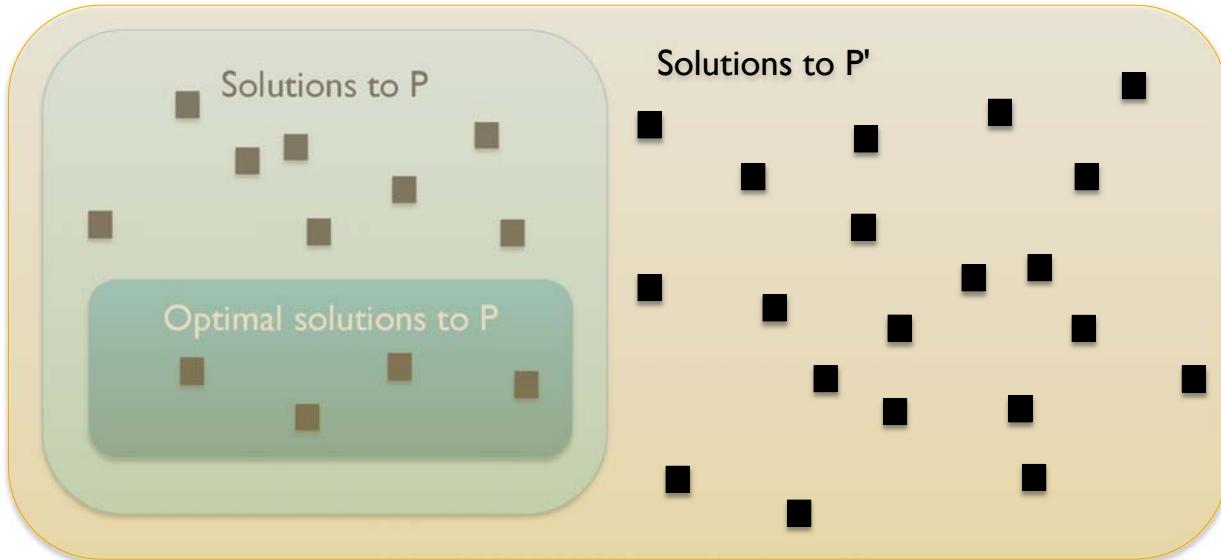
Includes the optimal solutions to  $P$ , so  $\min \{ \dots \}$  cannot be greater



# Fundamental ideas (3)



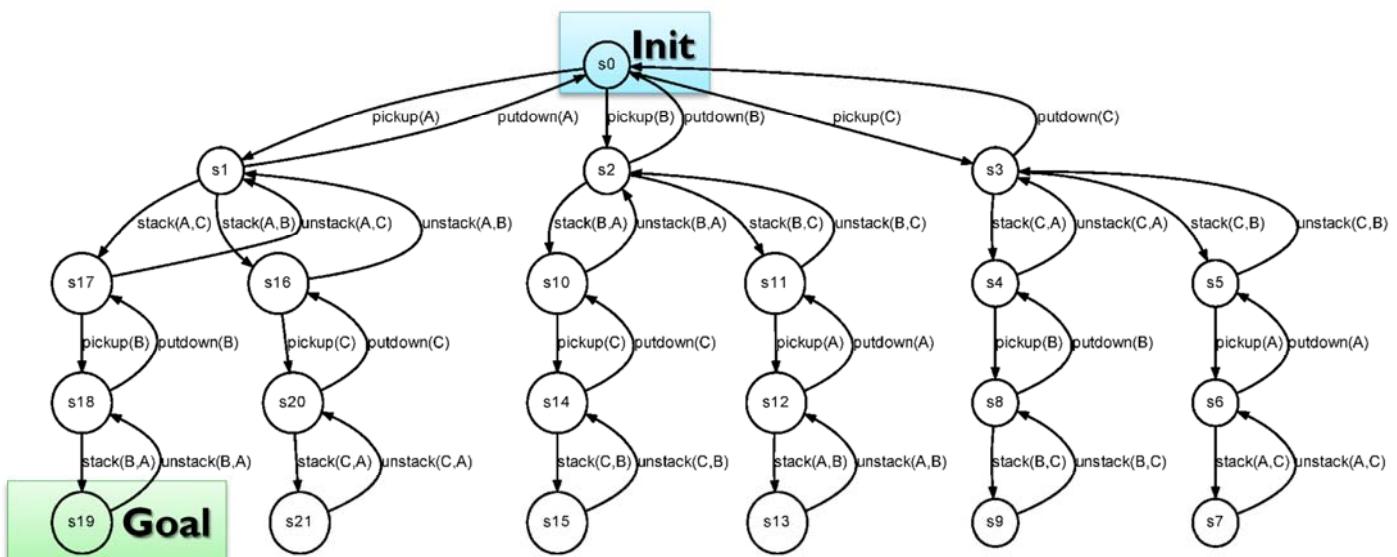
- Another sufficient criterion: All solutions to P remain solutions for P'
  - Stronger, but often easier to prove
- Relaxes the constraint on what is accepted as a solution:  
The *is-solution(plan)* test is "expanded, relaxed" to cover additional plans
  - Called relaxation



## Relaxation Example: Basis



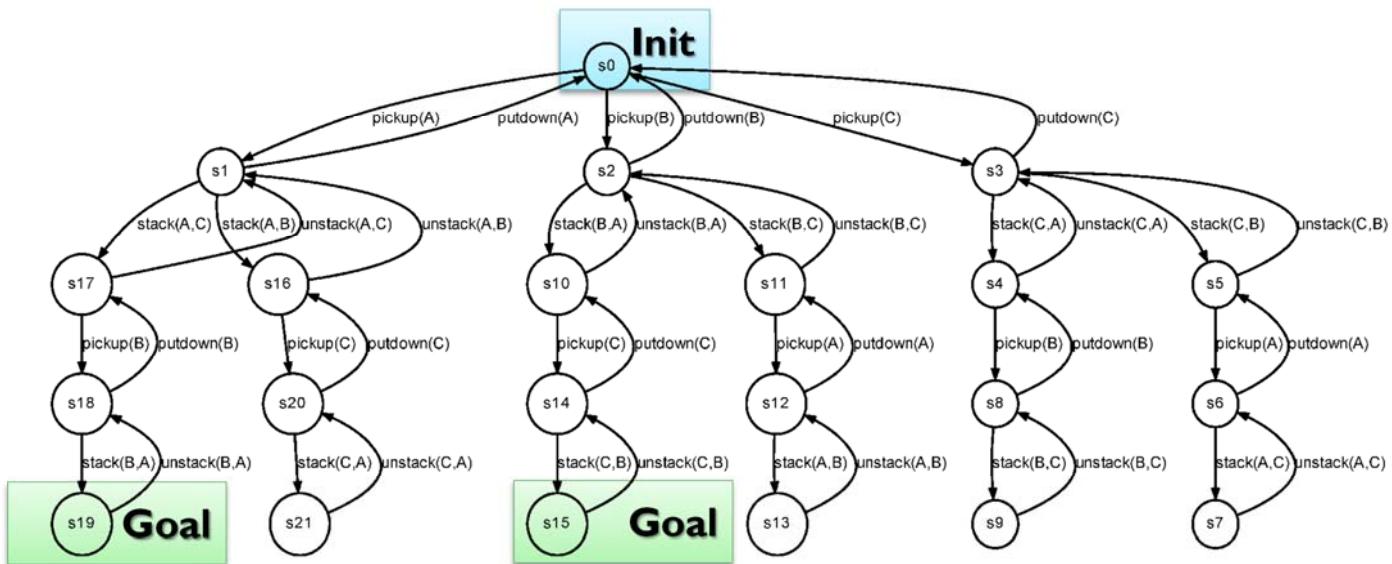
- A simple planning problem (domain + instance)
  - Blocks world, 3 blocks
  - Initially all blocks on the table
  - Goal: (and (on B A) (on A C)) (only satisfied in s19)
  - Solutions: All paths from init to goal (infinitely many – can have cycles)



# Relaxed Problem 1

## ■ Relaxation "method 1": Adding goal states

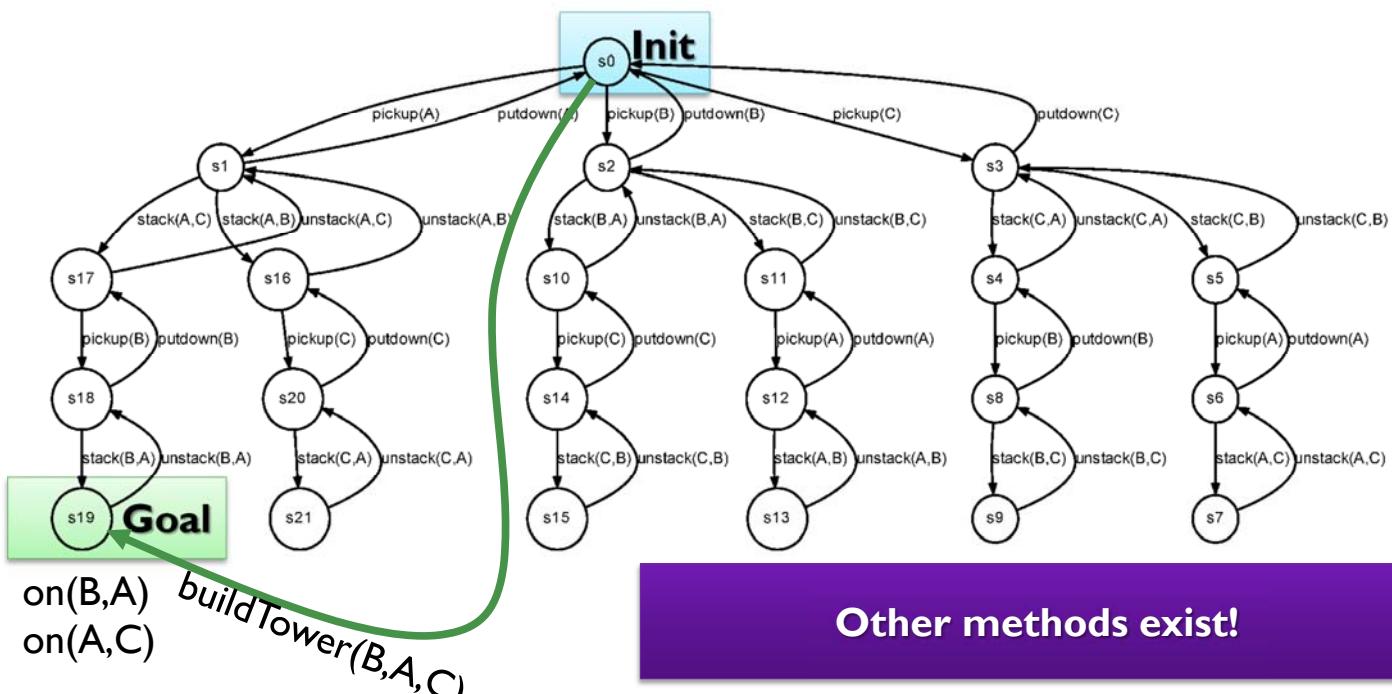
- New goal formula: (and (on B A) **(or (on A C) (on C B))**)
  - All old solutions still valid
  - Some non-solution paths become solutions



# Relaxed Problem 2

## ■ Relaxation "method 2": Adding new actions to the domain

- **Modifies** the state transition system (states/actions)
  - This particular example: a shorter solution becomes possible



# Understanding Relaxation



## ■ Important:

Relaxation is not simply "removing preconditions"

There are **many other relaxations**

Key: Preserving old solutions!

You **cannot** "use a relaxed problem as a heuristic".

What would that mean? The problem is not a number or function...

You use the **cost** of an **optimal solution** to the relaxed problem as a heuristic.

If you just take the cost of **any** solution to the relaxed problem,  
it can be inadmissible!

You have to solve it **optimally** to get the admissibility guarantee.

You don't just solve the relaxed problem once.

**Every time you reach a new state and want to calculate a heuristic value,**  
you have to solve the relaxed problem  
of getting from **that** state to the goal.

# Understanding Relaxation (2)



You've seen relaxations adapted to the 8-puzzle

What can you do automatically,  
independently of the planning domain used?

# Example: Pattern Database Heuristics

## PDB 1: State Variables



- Let's use a non-Boolean state variable representation
  - Example for 4 blocks:
    - **aboveA**  $\in \{ \text{clear}, \text{B}, \text{C}, \text{D}, \text{holding} \}$
    - **aboveB**  $\in \{ \text{clear}, \text{A}, \text{C}, \text{D}, \text{holding} \}$
    - **aboveC**  $\in \{ \text{clear}, \text{A}, \text{B}, \text{D}, \text{holding} \}$
    - **aboveD**  $\in \{ \text{clear}, \text{A}, \text{B}, \text{C}, \text{holding} \}$
    - **posA**  $\in \{ \text{on-table}, \text{other} \}$
    - **posB**  $\in \{ \text{on-table}, \text{other} \}$
    - **posC**  $\in \{ \text{on-table}, \text{other} \}$
    - **posD**  $\in \{ \text{on-table}, \text{other} \}$
    - **hand**  $\in \{ \text{empty}, \text{full} \}$

# PDB 2: Ignoring Variables



- Create an **abstraction** of the planning problem
  - **Ignoring** some of the state variables
  - Everywhere: In preconditions, effects, states, and goals
- Example: Care about { **aboveB**, **aboveD**, **posB**, **posD** }
  - **Rewrite** the goal
    - Original: { aboveA = B, aboveB = C, aboveC = D, aboveD = clear, hand = empty }
    - Abstract: { **aboveB = C**, **aboveD = clear** }
  - **Rewrite** actions, removing some preconditions / effects
    - (unstack D C) still requires aboveD = clear
    - But not necessarily aboveC = D
  - ...

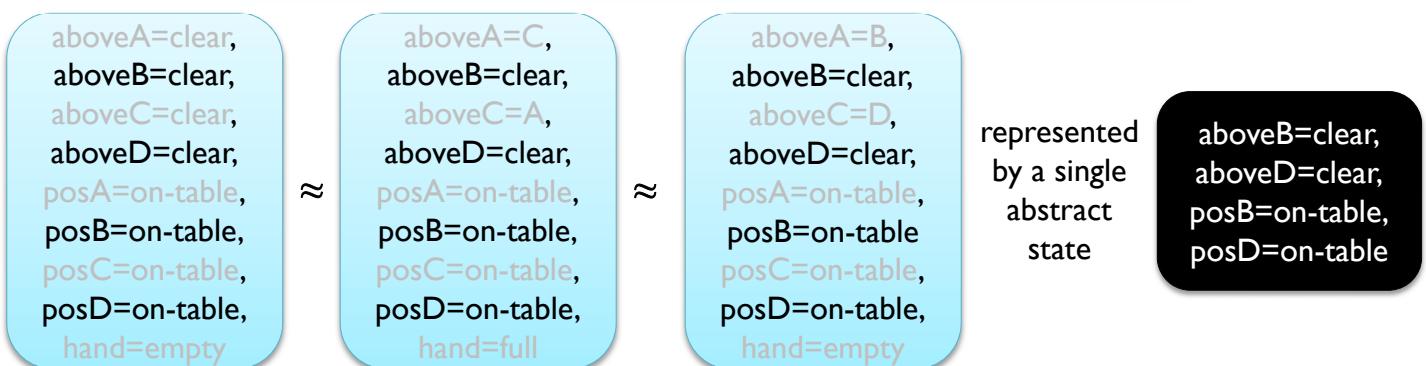
D  
C  
B  
A

# PDB 3: Patterns



- The **set of included variables** is called a **pattern**
  - States containing only these variables are called **abstract states**: Many states **match** a single abstract state

## Example: Pattern { **aboveB**, **aboveD**, **posB**, **posD** }



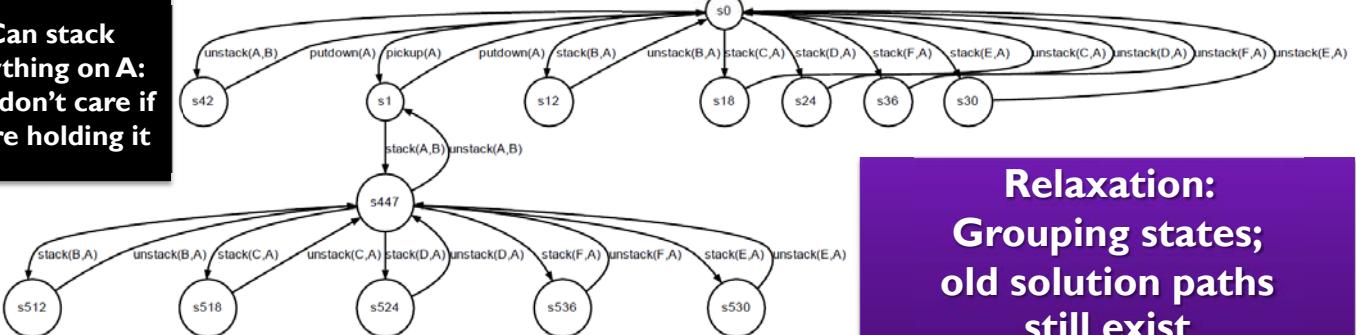
# PDB 4: State Space Size



- Abstract states reachable from "all on table", by pattern...

Blocks	All variables	Pattern={aboveA}	{aboveA,aboveB}
4	125	10	96
5	866	12	140
6	7057	14	192
7	65990	16	252
8	695417	18	320
9	8145730	20	396

Can stack anything on A:  
We don't care if we're holding it



Relaxation:  
Grouping states;  
old solution paths  
still exist

# PDB 5: Main Idea



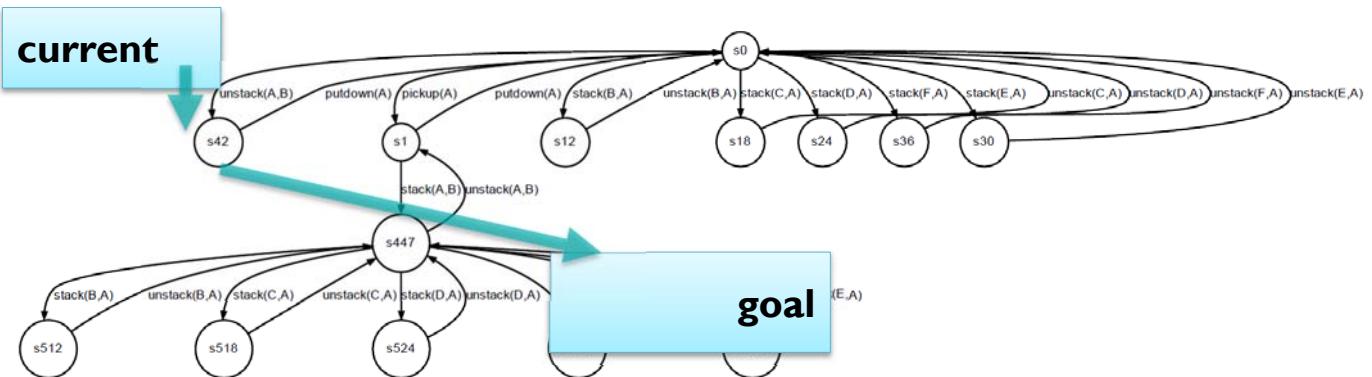
- To calculate  $h(s)$  for a newly encountered state  $s$ :

- Convert to abstract state

aboveA=B,  
aboveB=clear,  
aboveC=D,  
...

aboveA=B

- Find optimal path to abstract goal state – in a much smaller search space!
  - Relaxation → its cost is an admissible heuristic



# PDB 6: Main Idea

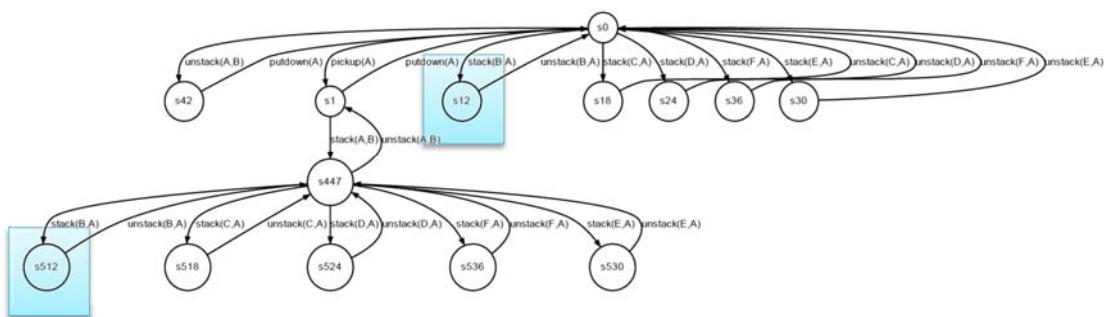


- To make PDB heuristics **more informative**:
  - Calculate costs for **several** patterns
    - {aboveA,aboveB} and {aboveC,aboveD} and ...
  - Take the **maximum** of the computed heuristic values
- For even more information:
  - Use patterns satisfying specific **additivity constraints**, measuring *different aspects* of the total cost
  - The **sum** of the computed heuristic values is admissible
- Real difficulty:
  - Choosing which patterns to use...

# PDB 7: Databases



- Where did the databases go?
  - Planning visits many **abstract** states over and over again
  - For each pattern, we **precompute** all reachable abstract states



- Determine which ones satisfy the **abstract goal**
  - Abstract: { aboveA = B }
- Calculate shortest paths from **every** abstract state to **any** abstract goal state
  - Dijkstra's algorithm backwards – possible in a small search space
  - Store in a fast look-up table, a *database*

# Example: Landmark Heuristics

Intention:

Get a rough idea about how planners might find guidance!

## Landmark Heuristics (1)



### Landmark:

"a geographic feature used by explorers and others  
to find their way back or through an area"



# Landmark Heuristics (2)



## Landmarks in planning:

Something you must pass by/through in every solution to a specific planning problem

Assume we are currently in state  $s$ ...

### Fact Landmark for $s$ :

A fact that is not true in  $s$ ,  
but must be true at some point  
in every solution starting in  $s$



clear(A)  
holding(C)

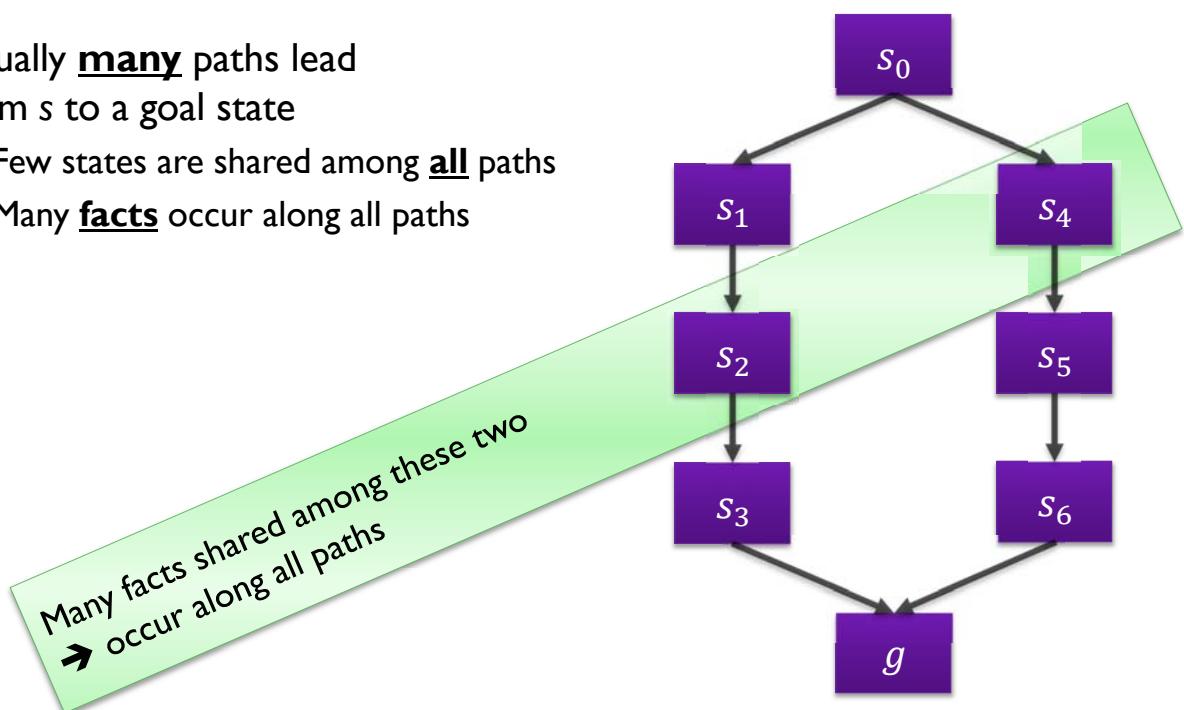
...

# Landmark Heuristics (3)



## Facts and formulas, not complete states!

- Usually many paths lead from  $s$  to a goal state
  - Few states are shared among all paths
  - Many facts occur along all paths



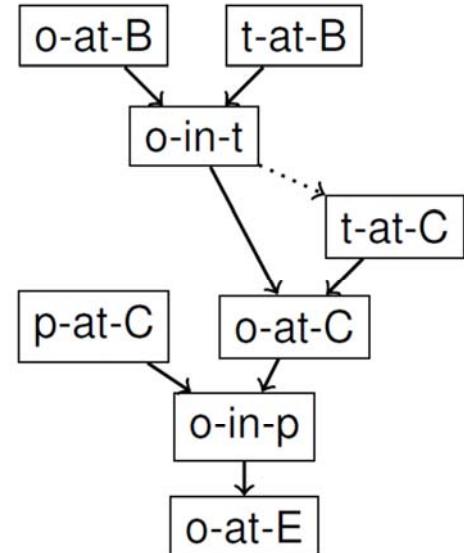
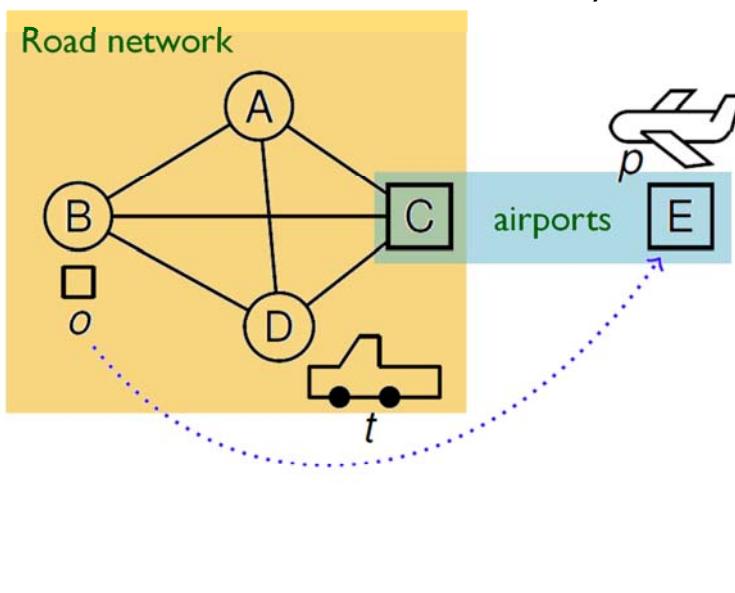
Finding landmarks: Complicated; not covered here

# Landmarks as Subgoals



## ■ Use of landmarks:

- As subgoals: Infer necessary orderings between landmarks, then try to achieve each landmark in succession
  - Example from Karpas & Richter:  
*Landmarks – Definitions, Discovery Methods and Uses*

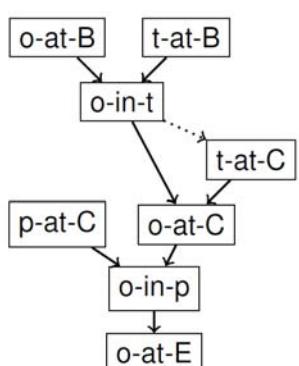


# Landmark Counts and Costs



## ■ Use of landmarks:

- As a basis for heuristic estimates
  - Used by LAMA, the winner of the sequential *satisficing* track of the International Planning Competition in 2008, 2011
- LAMA counts landmarks:
  - Identifies a set of landmarks that still need to be achieved after reaching state  $s$  through path (action sequence)  $\pi$
  - $L(s, \pi) =$   $(L \setminus \text{Accepted}(s, \pi)) \cup \text{ReqAgain}(s, \pi)$



$(L \setminus \text{Accepted}(s, \pi))$

All discovered landmarks, minus those that are accepted as achieved (has become true after predecessors are achieved!)

$\cup$   $\text{ReqAgain}(s, \pi)$

Plus those we can show will have to be re-achieved

# Search Algorithms for Non-Admissible Heuristics

## Search Algorithms for non-admissible heuristics



- Which search algorithm to use?
  - A\* expands many nodes (expensive),  
and for non-admissible heuristics, *there is still no optimality guarantee!*
  - Some planners use variations on **Hill Climbing**

**function** HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

*current*  $\leftarrow$  MAKE-NODE(*problem.INITIAL-STATE*)  
**loop do**

*neighbor*  $\leftarrow$  a highest-valued successor of *current*  
    **if** *neighbor.VALUE*  $\leq$  *current.VALUE* **then return** *current.STATE*  
    *current*  $\leftarrow$  *neighbor*

**Successors:** States created by applying an action to the current state

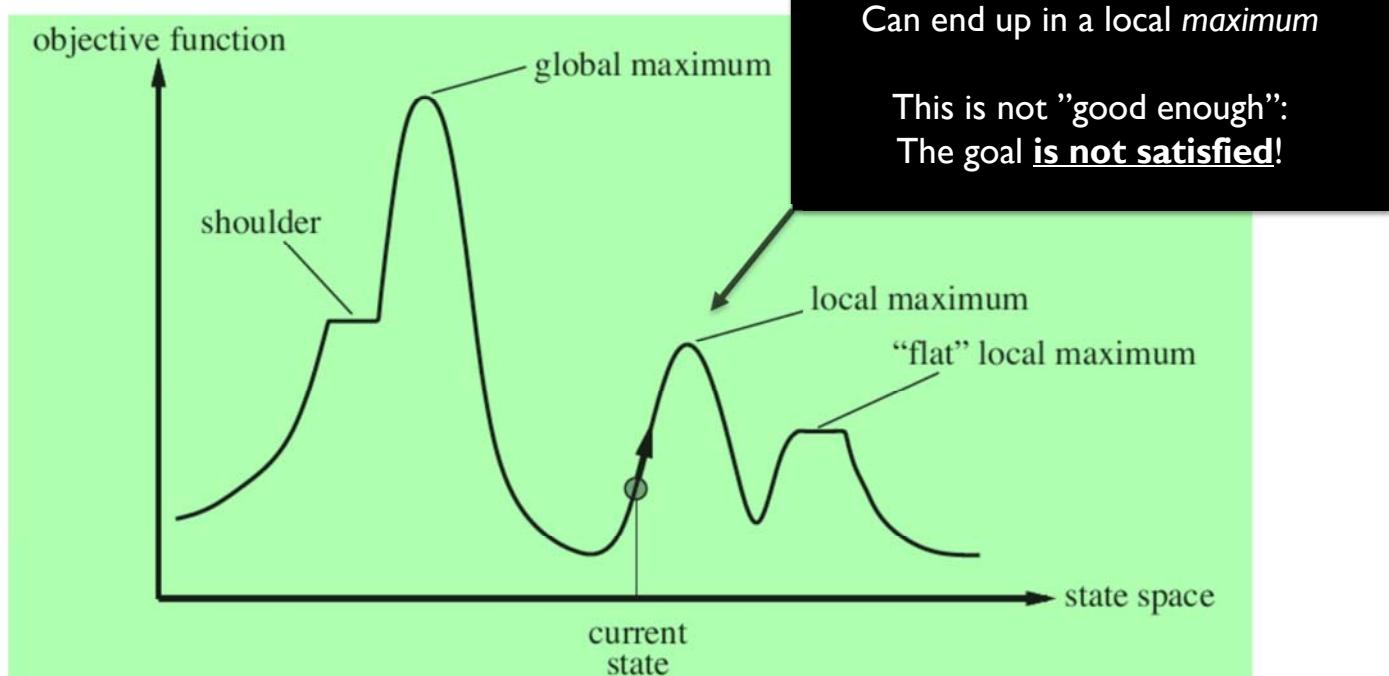
**Value:** Calculated as  $-h(s)$   
(maximize  $-h(s)$   $\rightarrow$  minimize  $h(s)$ )

Only considers children of the current node  
→ tries to move quickly towards low heuristic values  
→ doesn't examine "all nodes" with  $\text{cost}(n) + f(n) \leq \text{some value}$

# Local Optima in Planning

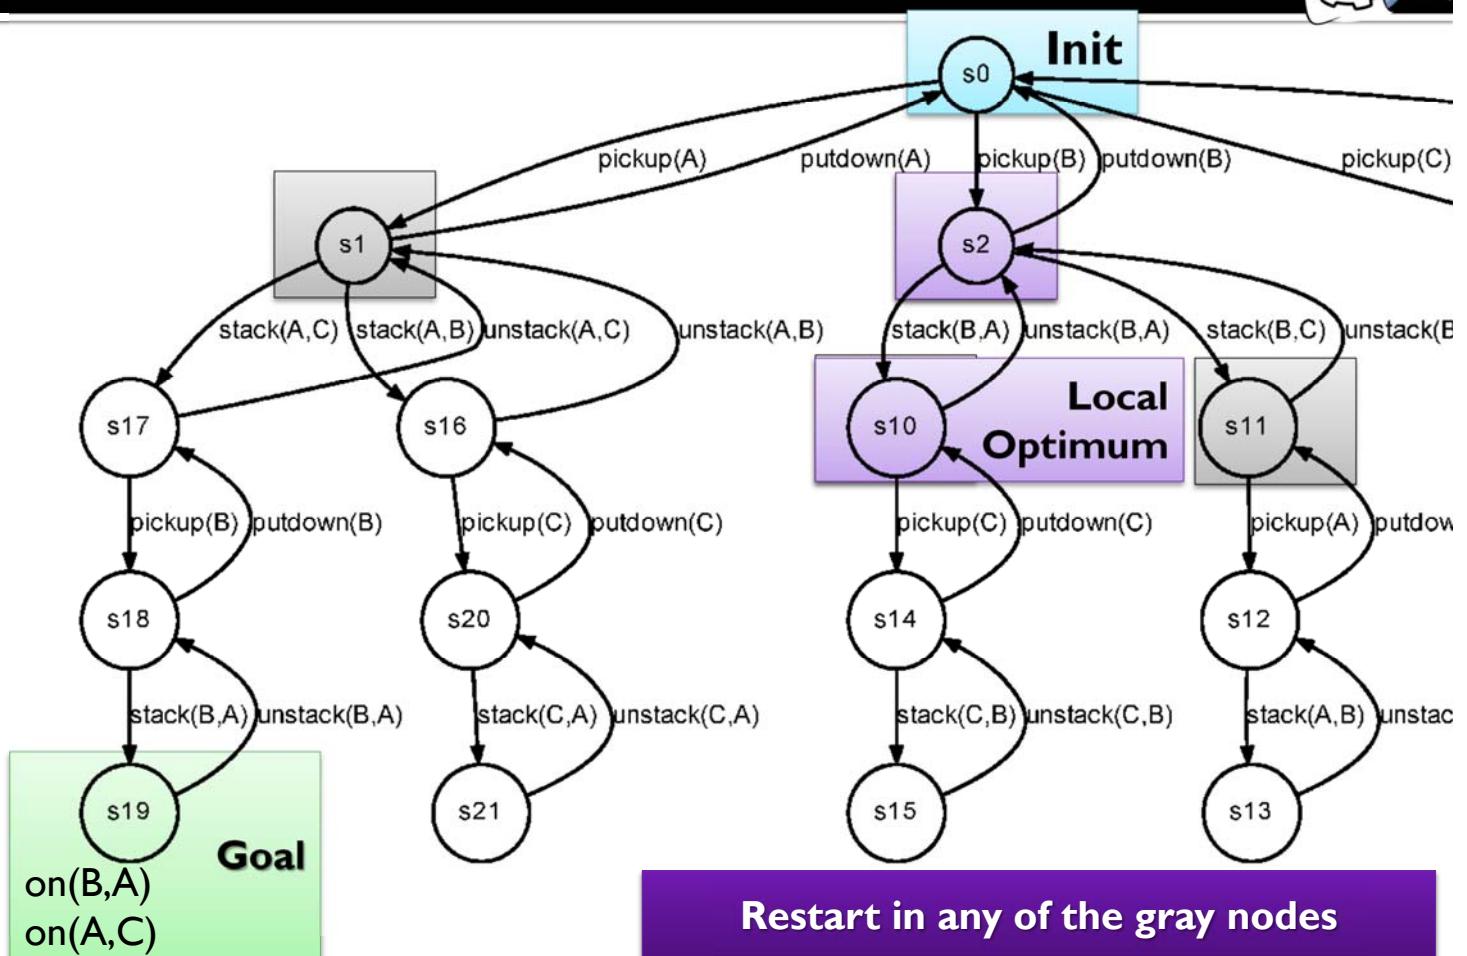


## Problem:



Possible solution:  
Restart search from another expanded state

# Local Optima: Restarting



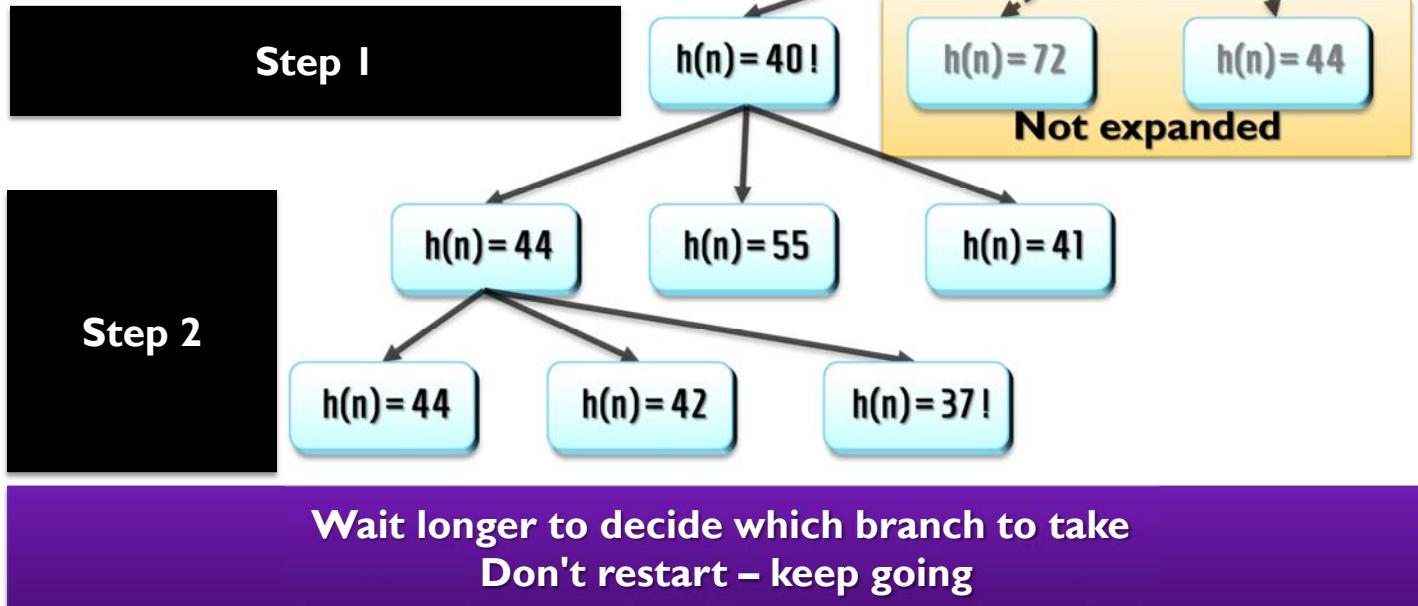
# Enforced Hill Climbing



- FastForward uses enforced hill climbing – approximately:

- $s \leftarrow \text{init-state}$
- repeat**

**expand** breadth-first until a better state  $s'$  is found  
**until** a goal state is found



## Summary of Forward Search



- Forward State Space Search:

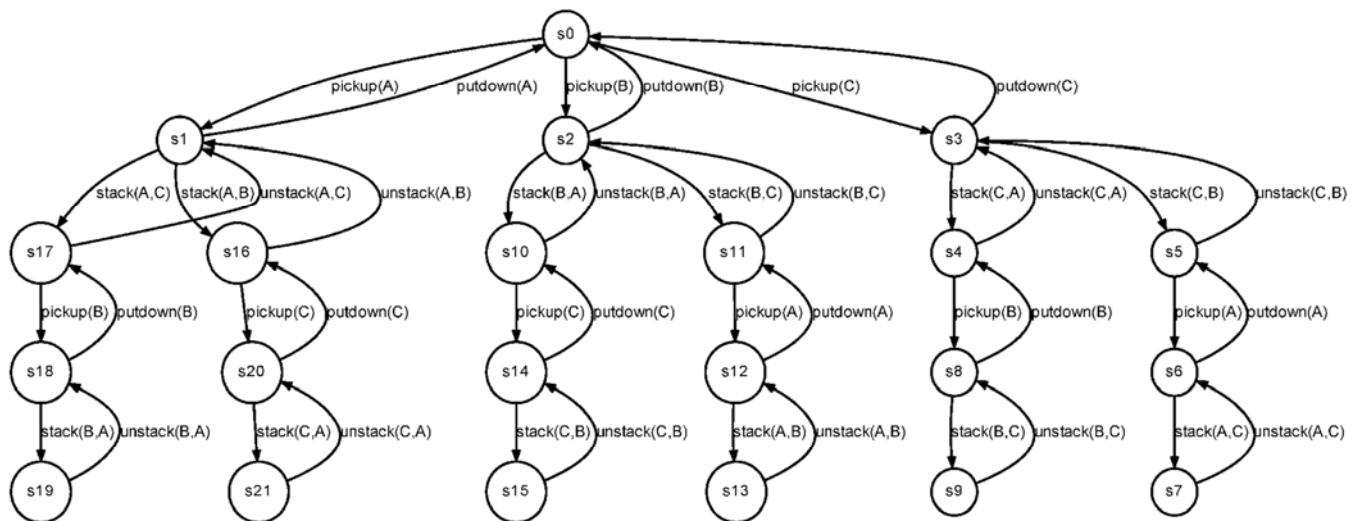
- Very common planning framework
  - Many** heuristic functions defined
    - TDDC17: Focus on understanding the basics, and general principles for defining general heuristics without knowing the problem to be solved (blocks, helicopters, ...)
    - TDDD48: Many additional heuristics (and principles for creating them)
  - Many** search strategies defined
    - TDDC17: Focus on applying standard strategies from general search
    - TDDD48: Additional planning-inspired strategies and their relations to specific heuristic functions

# Plan-Space Search: An overview

## Plan-Space Search



- Just because we are looking for a path in this graph...

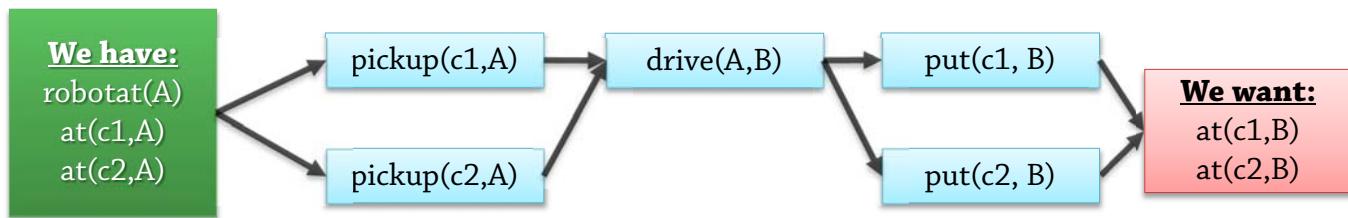


- ...doesn't mean we have to use it as a search space!

What if a search step could insert actions anywhere in a plan?

- Idea: Partial Order Causal Link (POCL) planning

- Means-ends analysis:
  - Find a condition (goal or precondition) that remains to be achieved
  - Add an action that achieves this condition
- But use a **partial order** for actions!
  - Insert actions "at any point" in a plan
  - Least/late commitment to ordering – "don't decide until you have to"



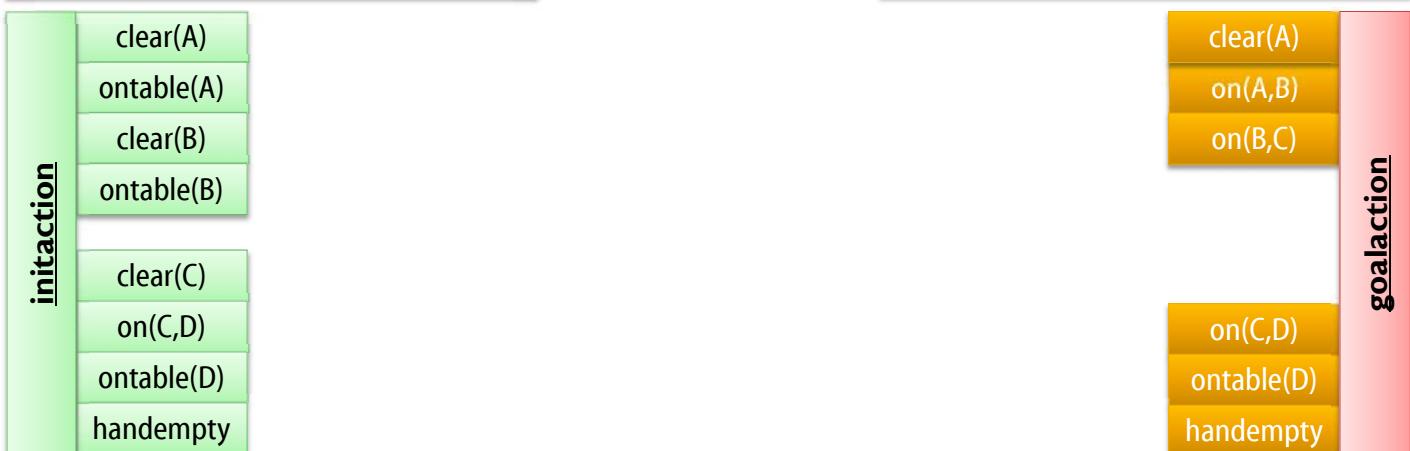
More sophisticated "bookkeeping" required!

## Partial-Order Planning

- How to keep track of initial state and goal?

**A "fake" initial action,  
whose effects are  
the facts of the initial  
state**

**A "fake" goal action,  
whose preconditions are  
the conjunctive goal facts**



Streamlines the planner: *Initial state and action effects are handled equivalently*  
*goal and action preconditions are handled equivalently.*

# Partial-Order Planning



- No sequence → must have explicit precedence constraints

A "fake" initial action,  
whose effects are  
the facts of the initial  
state

A "fake" goal action,  
whose preconditions are  
the conjunctive goal facts



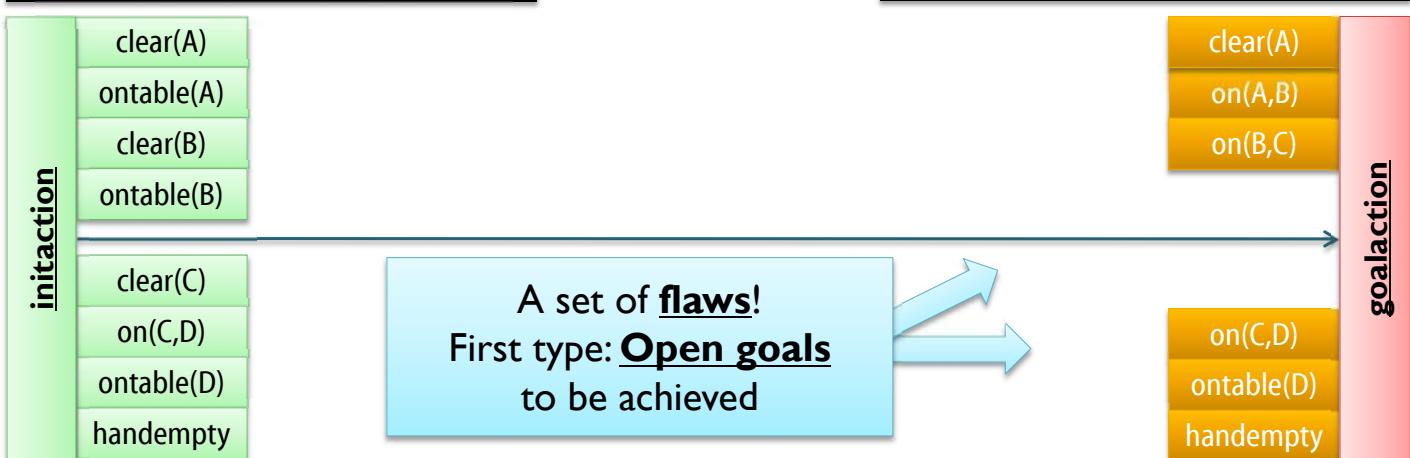
# Partial-Order Planning



- No current state – how to keep track of what remains to do?

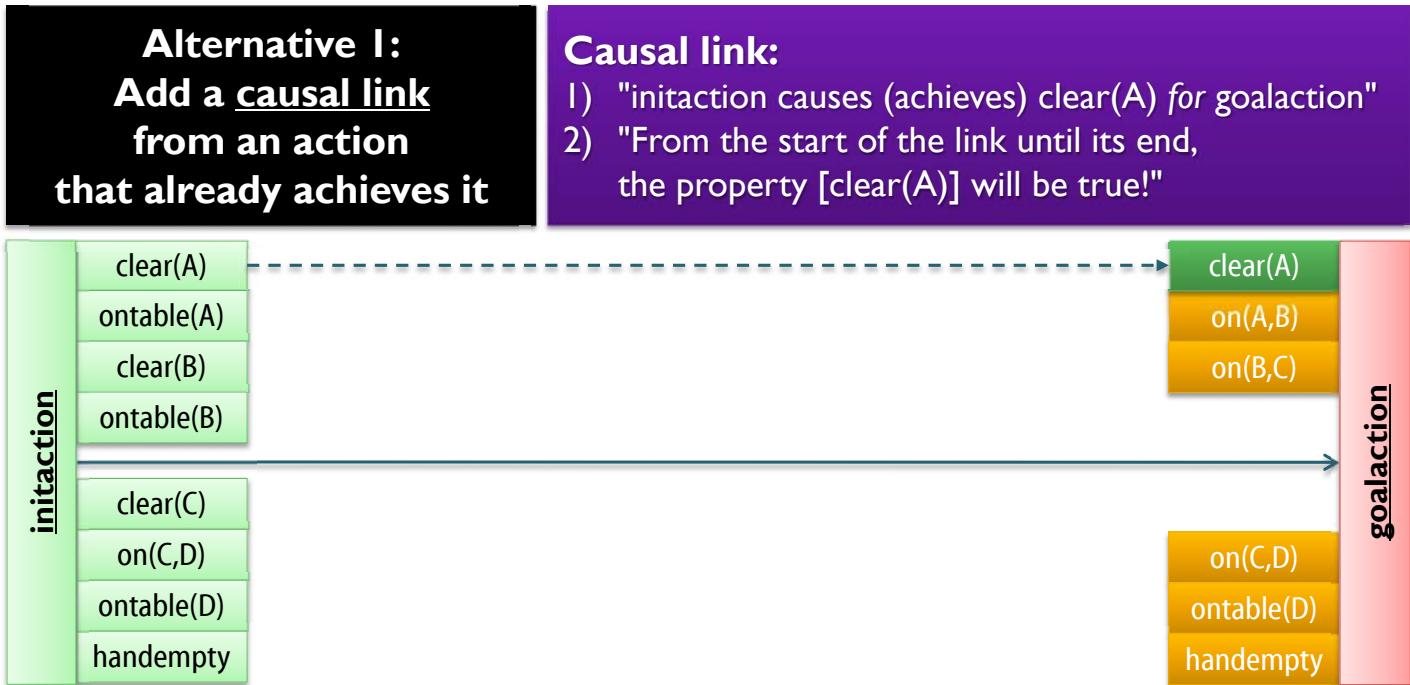
A "fake" initial action,  
whose effects are  
the facts of the initial  
state

A "fake" goal action,  
whose preconditions are  
the conjunctive goal facts



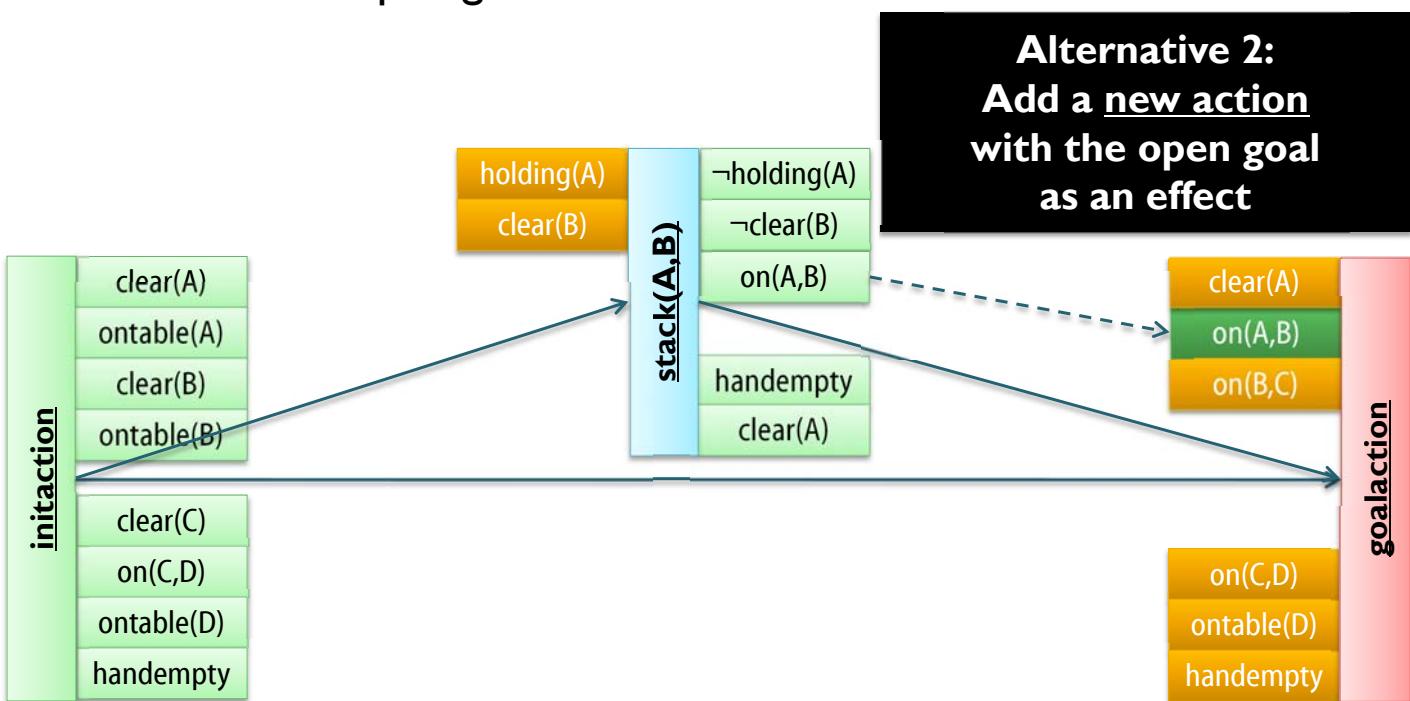
# Resolving Open Goals 1

- To resolve an open goal:



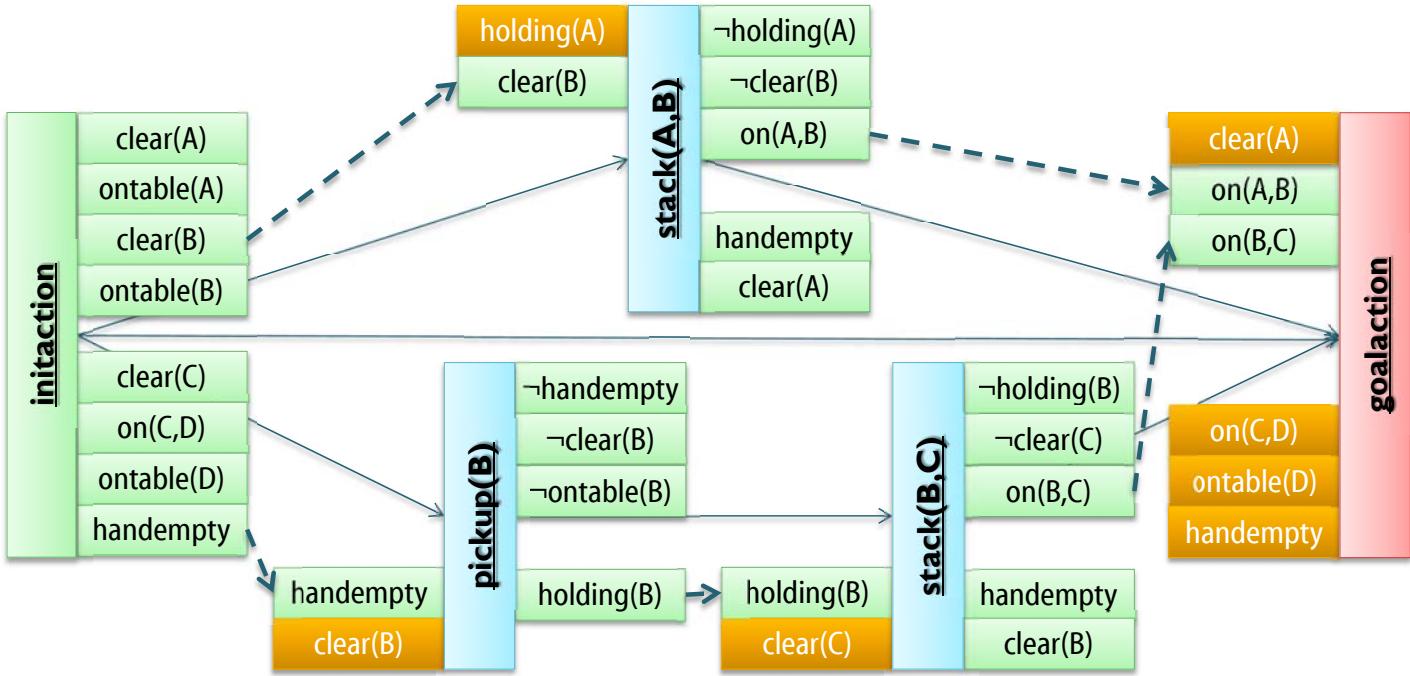
# Resolving Open Goals 2

- To resolve an open goal:



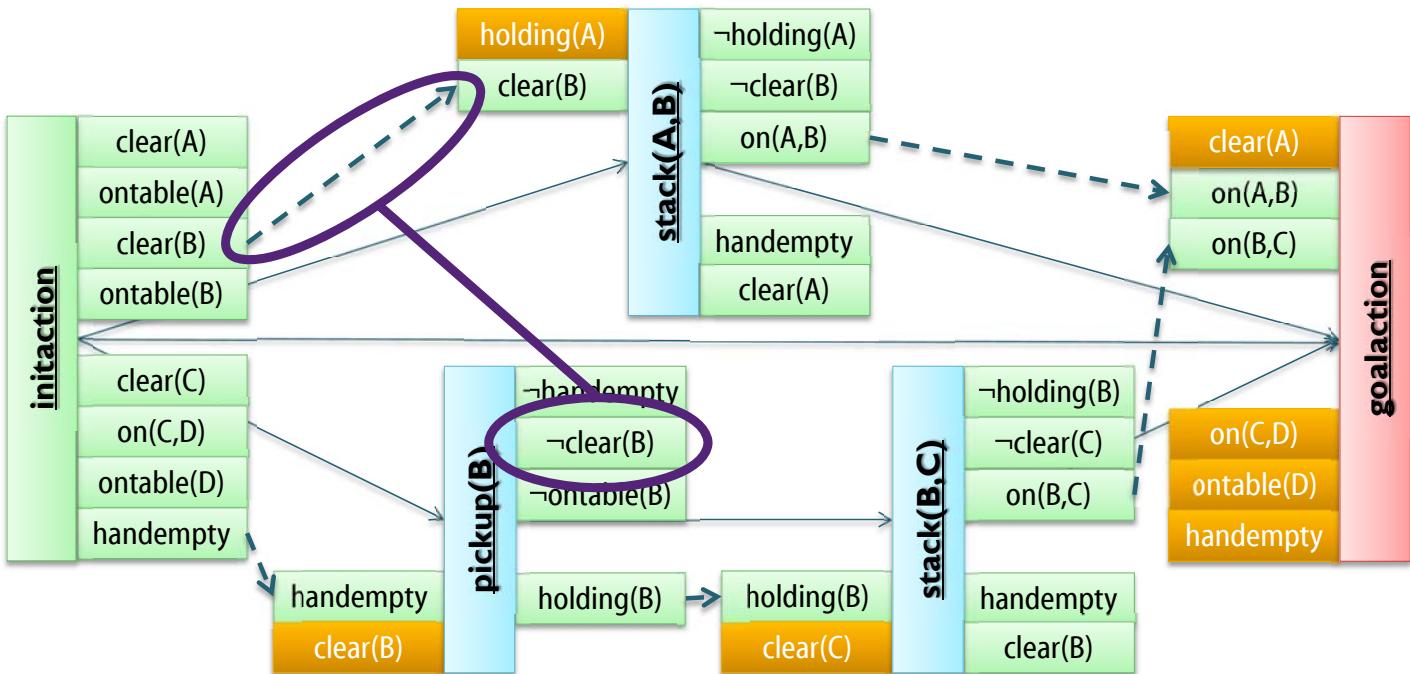
# Partial Orders

- Some pairs of actions can remain unordered



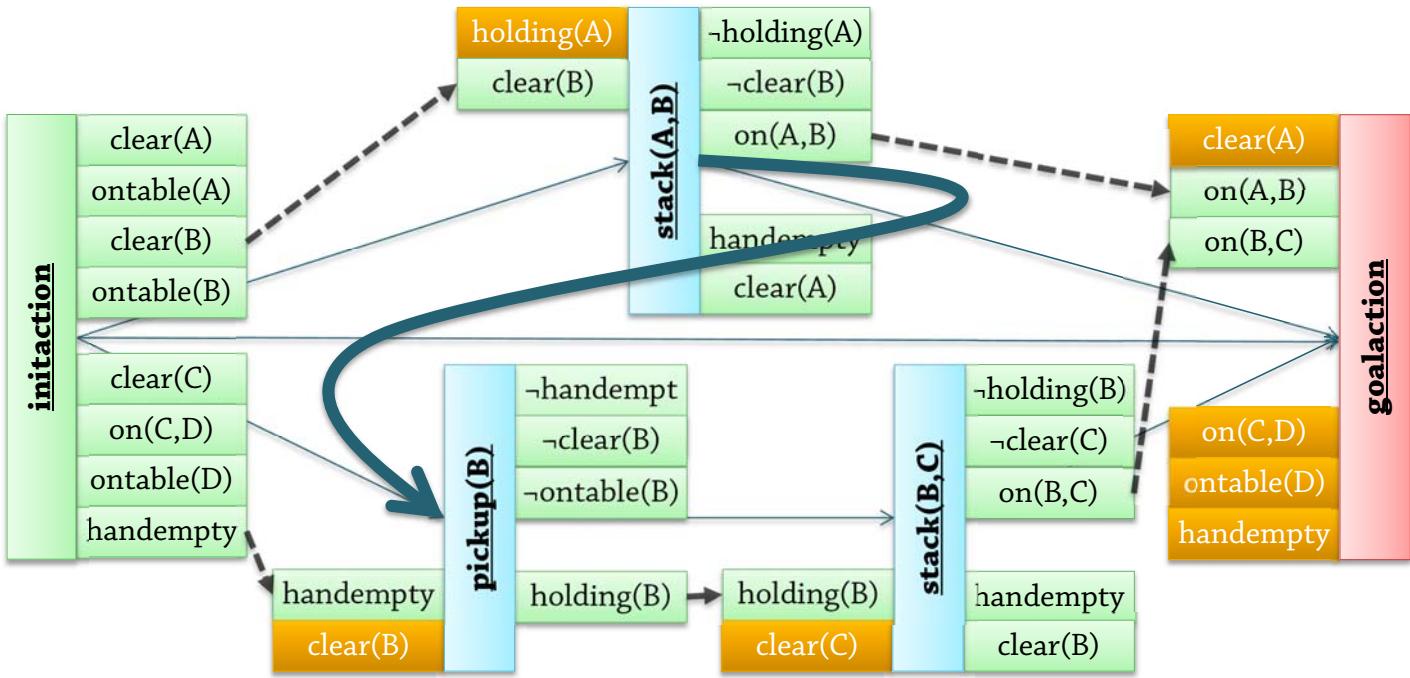
## Threats

- Second flaw type: A threat to be resolved
  - initaction supports clear(B) for stack(A,B) – there's a causal link
  - pickup(B)  deletes clear(B), and may occur "during" the link
  - We can't be certain that clear(B) still holds when stack(A,B) starts!



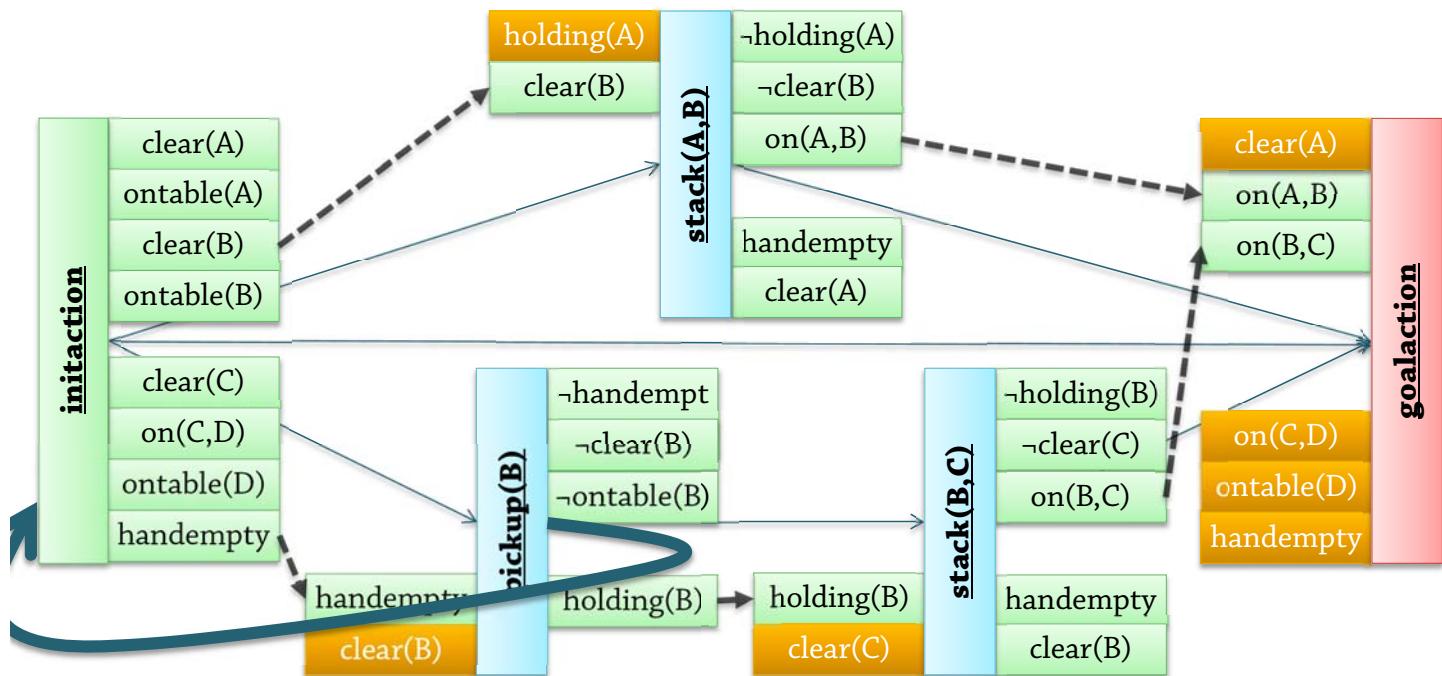
# Resolving Threats 1

- How to make sure that  $\text{clear}(B)$  holds when  $\text{stack}(A,B)$  starts?
  - Alternative 1: The action that disturbs the precondition is placed after the action that has the precondition
    - Only possible if the resulting partial order is consistent (acyclic)!



# Resolving Threats 2

- Alternative 2:
  - The action that disturbs the precondition is placed before the action that supports the precondition
  - Only possible if the resulting partial order is consistent – not in this case!



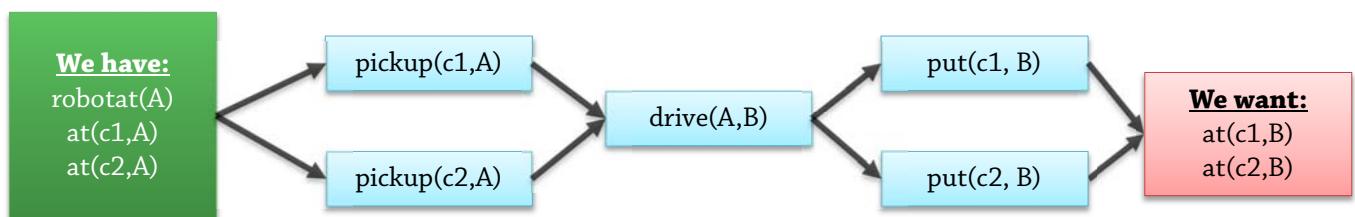
- Again, this is not sufficient in itself
  - Heuristics are required
- But it has advantages
  - No need to commit to order

**Remember what a flaw is!**

- Not something "wrong" in a final solution
- Not a mistake
- Simply bookkeeping for "something we haven't taken care of yet"
  - Open goal: Must decide how to achieve a precondition
  - Threat: Must decide how to order actions

## Partial-Order Solutions

- Original motivation: performance
  - Therefore, a partial-order plan is a solution iff all sequential plans satisfying the ordering are solutions
    - Similarly, executable iff corresponding sequential plans are executable
    - <pickup(c1,A), pickup(c2,A), drive(A,B), put(c1,B), put(c2,B)>
    - <pickup(c2,A), pickup(c1,A), drive(A,B), put(c1,B), put(c2,B)>
    - <pickup(c1,A), pickup(c2,A), drive(A,B), put(c2,B), put(c1,B)>
    - <pickup(c2,A), pickup(c1,A), drive(A,B), put(c2,B), put(c1,B)>

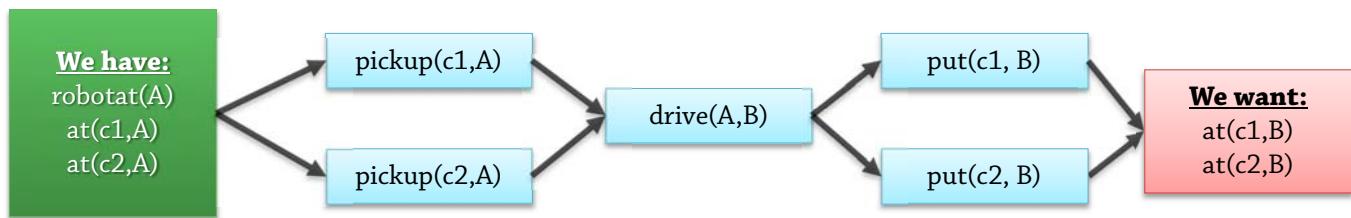


# Partial Orders and Concurrency



- Can be extended to allow concurrent execution

- Requires an extended action model!
  - Our model *does not define* what happens if  $c_1$  and  $c_2$  are picked up **simultaneously** (or even if we are allowed to do this)!



## Beyond State Space Planning

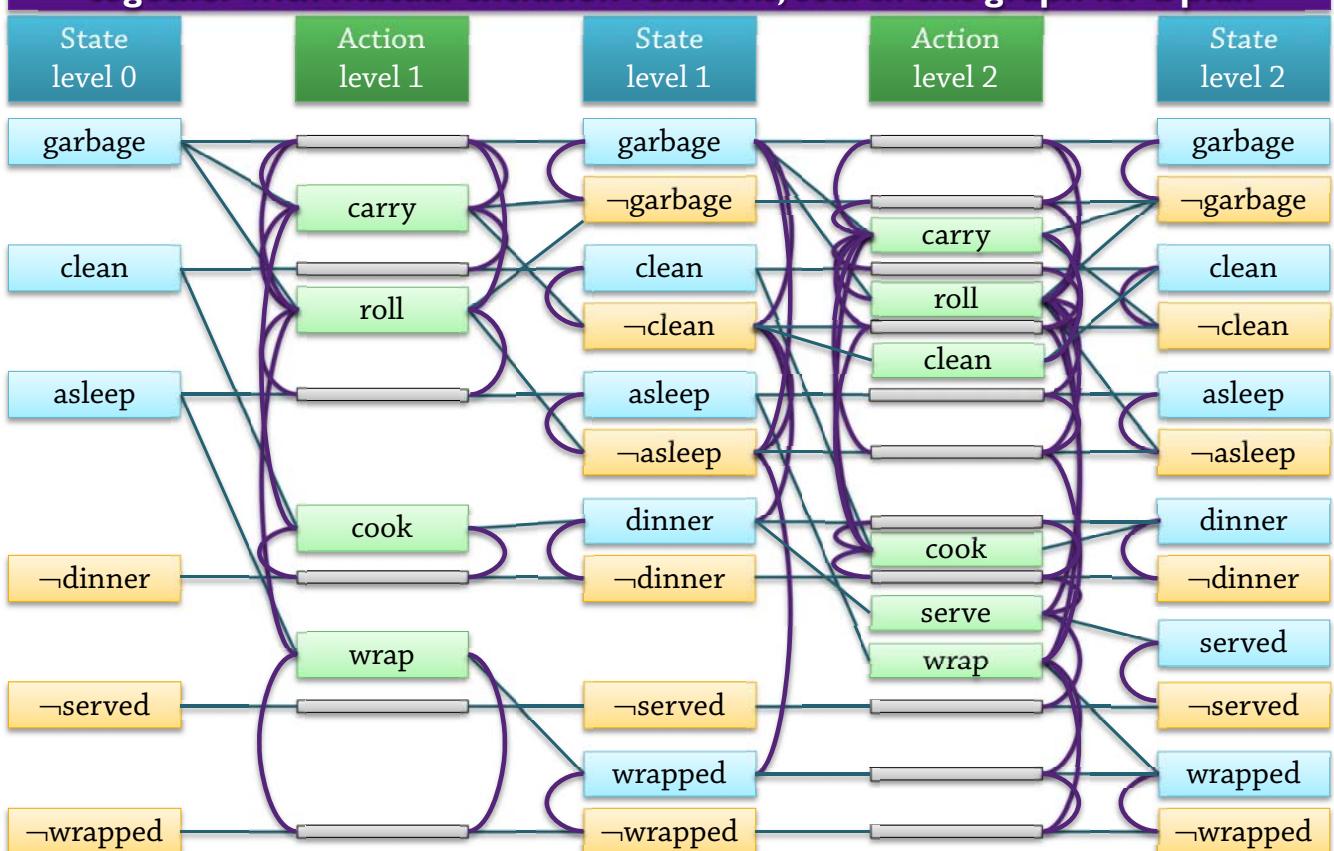
Many alternative search spaces exist;  
some examples here!

- SAT: Boolean satisfiability, first problem proven NP-complete
  - Translate planning to a SAT problem – propositions + complex formula
  - Apply existing SAT solver
  - Each **SAT solution** corresponds to a **solution plan**

"Action propositions": Which actions are executed, and when?							"Fact propositions": Which facts are true, and when?	
							$\varphi$	Solutions
...	...	...	...	...	...	...	$\varphi$	
-	-	-	-	-	-	-	-	
-	-	true	true	-	-	-	-	
-	true	-	true	-	true	-	-	
Seq. plan!		-	true	true	-	-	true	<b>true</b>
		true	-	-	true	true	-	-
		true	-	true	true	true	-	-
Seq. plan!		true	true	-	-	true	true	<b>true</b>
		true	true	true	true	true	true	-

## Planning Graphs

Compute a graph representing what might be achievable per time step, together with mutual exclusion relations; search this graph for a plan



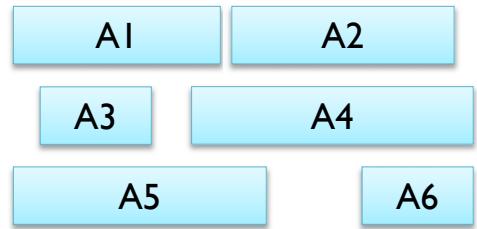
# Beyond Classical Planning: Time, Concurrency, Resources

## Time, Concurrency and Resources



- Many planners support time and concurrency
  - Actions have durations
  - Actions can be scheduled in parallel, suitable for multiple agents
    - Example: Duration determined by current state  
(:durative-action heat-water
      - :parameters (?p - pan)
      - :duration (= ?duration (/ (- 100 (temperature ?p)) (heat-rate)))
      - :condition (and (at start (full ?p))  
                 (at start (onHeatSource ?p))  
                 (at start (byPan))  
                 (over all (full ?p))  
                 (over all (onHeatSource ?p))  
                 (over all (heating ?p))  
                 (at end (byPan)))
      - :effect (and (at start (heating ?p))  
                 (at end (not (heating ?p))))  
                 (at end (assign (temperature ?p) 100)))

Conditions / effects  
at the *start* or *end*  
of the action



Implemented by quite a few planners –  
affects algorithms and heuristics!

# Time, Concurrency and Resources (2)



- Example: Duration selected by planner

(:durative-action heat-water

:parameters (?p - pan)

:duration (at end (<= ?duration (/ (- 100 (temperature ?p)) (heat-rate))))

:condition (and (at start (full ?p))

(at start (onHeatSource ?p))

(at start (byPan))

(over all (full ?p))

(over all (onHeatSource ?p))

(over all (heating ?p))

(at end (byPan)))

:effect (and (at start (heating ?p))

(at end (not (heating ?p))))

(at end (increase (temperature ?p) (\* ?duration (heat-rate))))))

General increase and decrease effects

for numeric values

(handles concurrent production and consumption of resources)

# Time, Concurrency and Resources (3)



- Some planners support continuous linear effects

- Occurring throughout the duration of the action

- (:durative-action fly

:parameters (?p - airplane ?a ?b - airport)

:duration (= ?duration (flight-time ?a ?b))

:condition (and (at start (at ?p ?a))

(over all (inflight ?p))

(over all (>= (fuel-level ?p) 0))))

:effect (and (at start (not (at ?p ?a)))

(at start (inflight ?p))

(at end (not (inflight ?p))))

(at end (at ?p ?b))

(decrease (fuel-level ?p) (\* #t (fuel-consumption-rate ?p))))))

- (:action midair-refuel

:parameters (?p)

:precondition (inflight ?p)

:effect (assign (fuel-level ?p) (fuel-capacity ?p)))

# Time, Concurrency and Resources (4)

## ■ Hybrid discrete/continuous planning

- (:durative-action navigate

:control-variables ((velX) (velY))

:duration (and (<= ?duration 5000))

:condition (and

(over all (>= (velX) -4)) (over all (<= (velX) 4))

(over all (>= (velY) -4)) (over all (<= (velY) 4))

(over all (<= (x) 700)) (over all (>= (x) 0))

(over all (<= (y) 700)) (over all (>= (y) 0))

(at start (AUV-ready)))

:effect (and

(at start (not (AUV-ready))))

(at end (AUV-ready))

(increase (x) (\* 1.0 (velX) #t))

(increase (y) (\* 1.0 (velY) #t))))

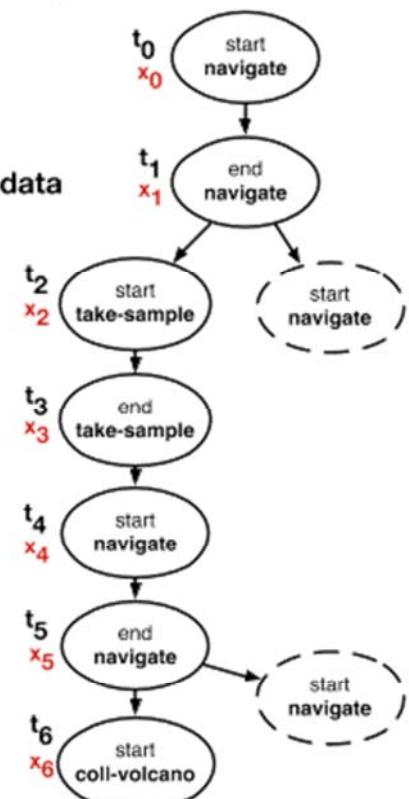
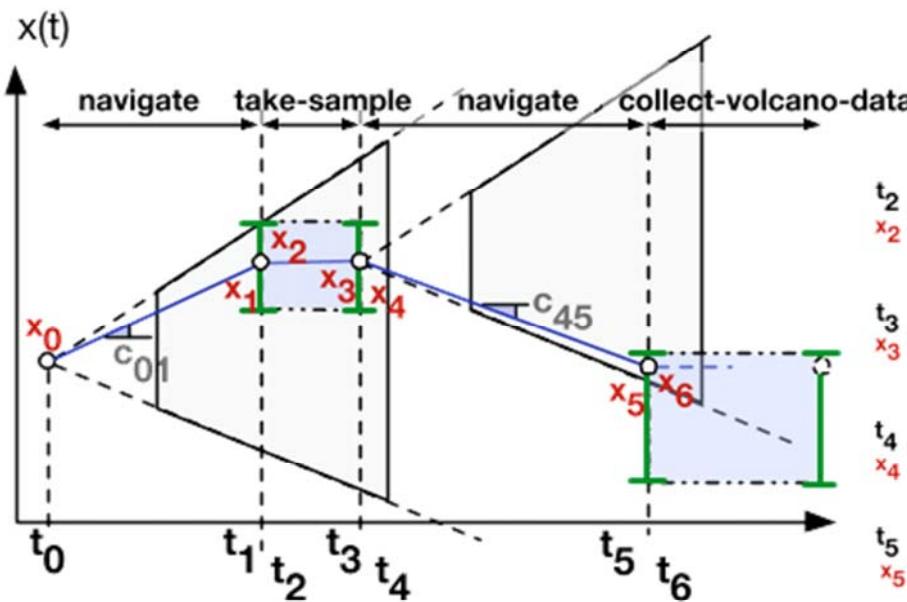
Can vary continuously  
during the action

Continuous effects  
depend on the  
**instantaneous** values of  
control variables

# Time, Concurrency and Resources (5)

## ■ Flow tubes result in reachable regions (shaded)

- Rectangles are conditions for **take-sample**, **collect**



# Beyond Classical Planning: Incomplete Information

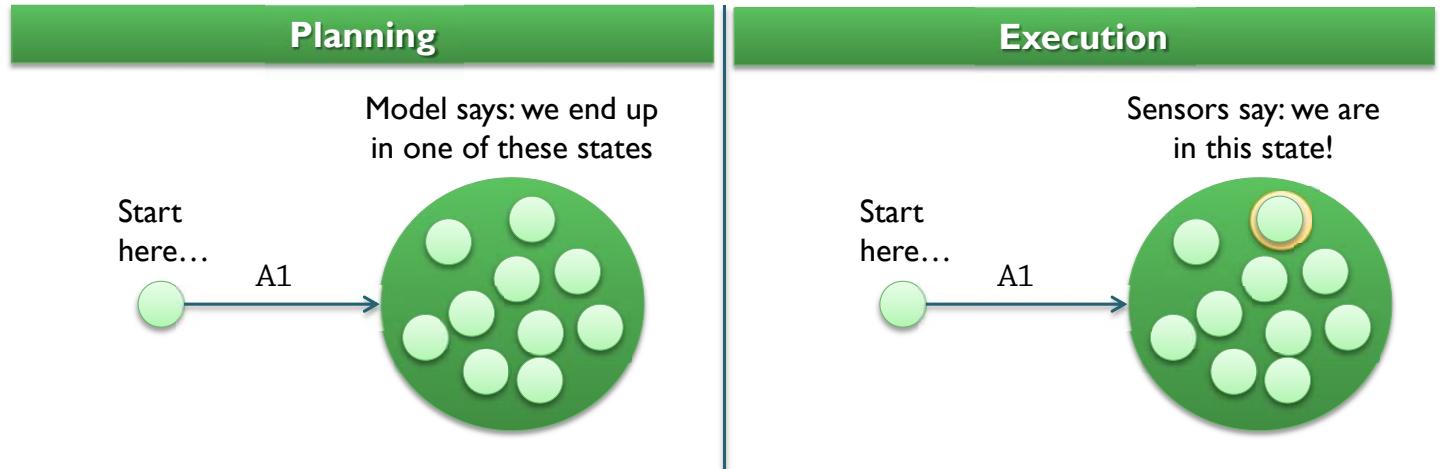
## Multiple Outcomes



- In reality, actions may have multiple outcomes
  - Some outcomes can indicate faulty / imperfect execution
    - **pick-up(object)**  
Intended outcome: carrying(object) is true  
Unintended outcome: carrying(object) is false
    - **move(100,100)**  
Intended outcome: xpos(robot)=100  
Unintended outcome: xpos(robot) != 100
    - **jump-with-parachute**  
Intended outcome: alive is true  
Unintended outcome: alive is false
  - Some outcomes are more random, but clearly desirable / undesirable
    - Pick a present at random – do I get the one I longed for?
    - Toss a coin – do I win?
  - Sometimes we have no clear idea what is desirable
    - Outcome will affect how we can continue, but in less predictable ways

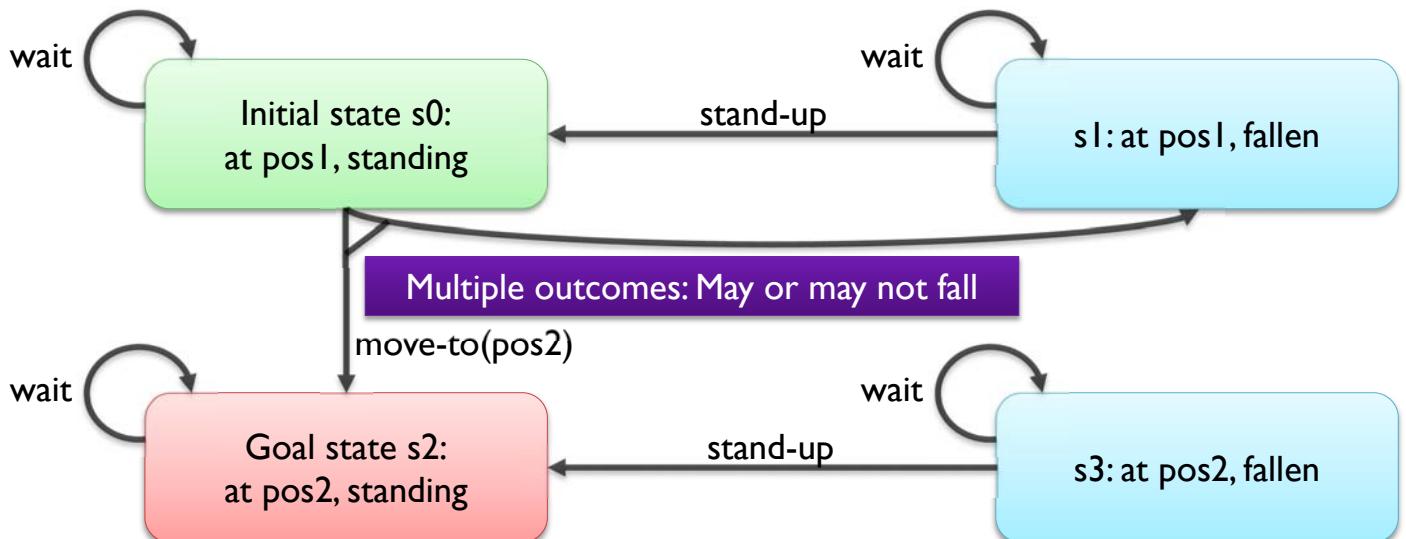
To a *planner*,  
there is generally  
no difference  
between these  
cases!

- **Nondeterministic planning:**
  - Many possible outcomes for each action
- FOND: **Fully Observable Non-Deterministic**
  - After executing an action, sensors determine exactly which state we are in



## FOND Planning: Plan Structure (1)

- Example state transition system:



- **Intuitive strategy:**
  - **while** (not in s2) {
    - move-to(pos2);**
    - if** (fallen) **stand-up;**
}

In FOND planning, action choice should depend on actual outcomes (through the current state)

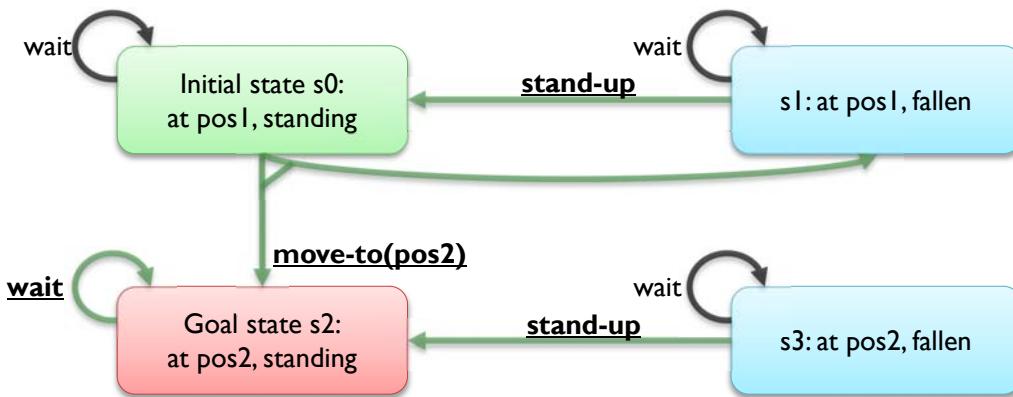
There may be no upper bound on how many actions we may have to execute!

# FOND Planning: Plan Structure (2)



- Examples of formal plan structures:

- **Conditional plans** (with if/then/else statements)
  - **Policies**  $\pi : S \rightarrow A$ 
    - Defining, for each state, which action to execute whenever we end up there
    - $\pi(s_0) = \text{move-to(pos2)}$
    - $\pi(s_1) = \text{stand-up}$
    - $\pi(s_2) = \text{wait}$
    - $\pi(s_3) = \text{stand-up}$
- Or at least, for every state  
that is *reachable* from the given initial state  
(A policy can be a *partial function*)



## Stochastic Systems

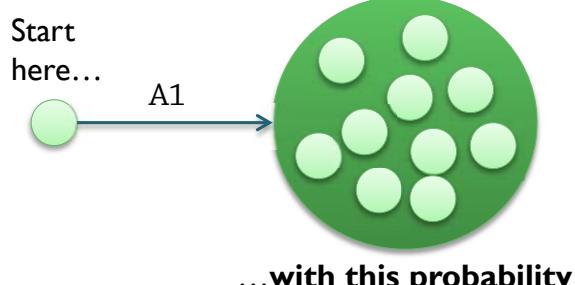


- **Probabilistic planning** → a stochastic system

- $P(s, a, s')$ : Given that we are in  $s$  and execute  $a$ , the probability of ending up in  $s'$
- For every state  $s$  and action  $a$ , we have  $\sum_{s' \in S} P(s, a, s') = 1$ :  
The world gives us 100% probability of ending up in *some* state

### Planning

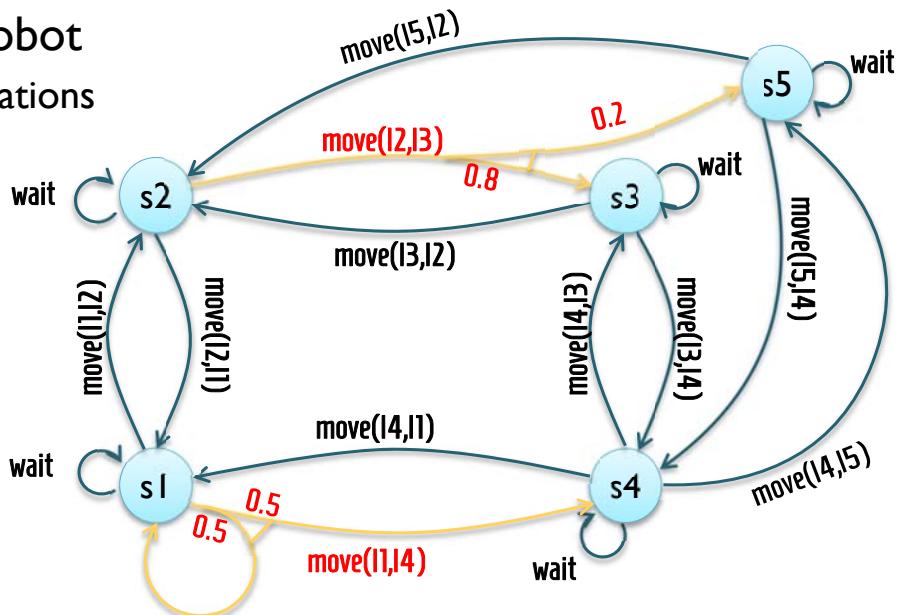
Model says: we end up in one of these states



# Stochastic System Example

- ## ■ Example: A single robot

- Moving between 5 locations
  - For simplicity,  
states correspond  
directly to  
locations
    - s1: at(r1, l1)
    - s2: at(r1, l2)
    - s3: at(r1, l3)
    - s4: at(r1, l4)
    - s5: at(r1, l5)



- Some transitions are deterministic, some are stochastic
    - Trying to move from 12 to 13: You may end up at 15 instead (20% risk)
    - Trying to move from 11 to 14: You may stay where you are instead (50% risk)
  - (Can't always move in both directions, e.g. due to terrain gradient)

# Representation Example: RDDL

- ## ▪ Relational Dynamic Influence Diagram Language – a different view

- Based on Dynamic Bayesian Networks

- **domain** prop\_dbn {  
    requirements = { reward - deterministic };  
    *// Define the state and action variables (not parameterized here)*

```
pvariables {  
    p : { state - fluent , bool , default = false };  
    q : { state - fluent , bool , default = false };  
    r : { state - fluent , bool , default = false };  
    a : { action - fluent , bool , default = false };  
};
```

// Define the conditional probability function for each next state variable in terms of previous state and action

**cpfs** {

$p' = \text{if } (p \wedge r) \text{ then Bernoulli (.9) else Bernoulli (.3);}$

$q' = \text{if } (q \wedge r) \text{ then Bernoulli (.9)}$

else if (a) then Bernoulli (.3) else Bernoulli (.8);

$r' = \text{if } (\sim q) \text{ then } \text{KronDelta}(r) \text{ else } \text{KronDelta}(r \Leftrightarrow q);$

};

*// Define the reward function ; note that boolean functions are  
// treated as 0/1 integers in arithmetic expressions*

$$\text{reward} = p + q - r;$$

}

	<u>Non-Observable:</u> No information gained after action	<u>Fully Observable:</u> Exact outcome known after action
<u>Deterministic:</u> Exact outcome known in advance	Classical planning (possibly with extensions) (Information dimension is irrelevant; we know everything before execution)	
<u>Non-deterministic:</u> Multiple outcomes, no probabilities	NOND: Conformant Planning	FOND: Conditional / Contingent Planning
<u>Probabilistic:</u> Multiple outcomes with probabilities	Probabilistic Conformant Planning  (Special case of POMDPs, Partially Observable MDPs)	Probabilistic Conditional Planning  Stochastic Shortest Path Problems  Markov Decision Processes (MDPs)

## Conclusions

# Example Questions



## ■ Example questions:

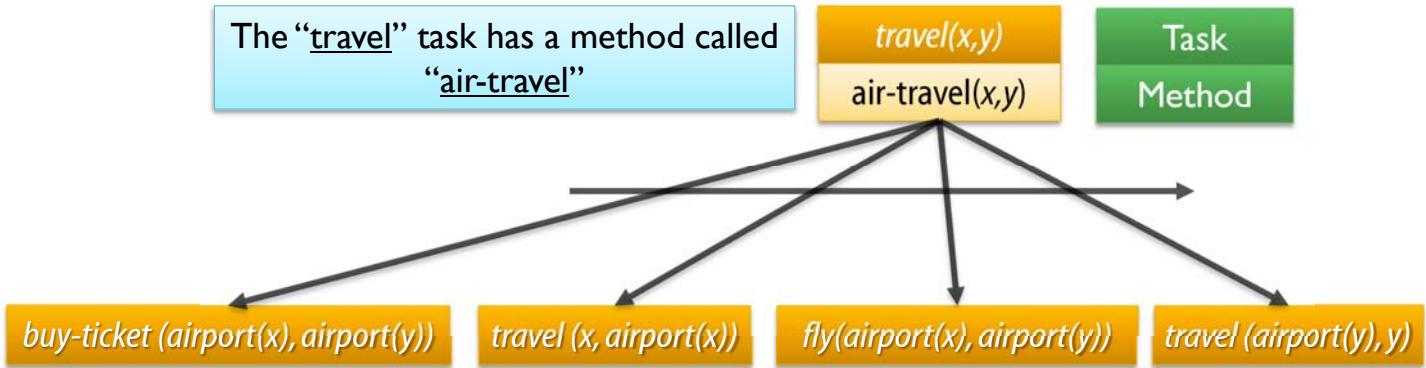
- Describe and explain:
  - Relaxation, and how it can be applied to create an admissible heuristic
  - PDB heuristics
  - Landmarks
  - The main advantage of partial-order planning over forward state space planning
  - The different kinds of flaws that can occur in a partially ordered plan
  - The ways in which you can resolve an open goal
  - Why planners don't use Dijkstra's algorithm to search the state space
- Is the number of open goals an *admissible* heuristic? Why / why not?
- Given a domain, problem and initial state:
  - Expand the forward state space by two levels
- How can you apply hill climbing to planning?
  - How do you have to modify standard hill climbing, and why?

# TDDD48 Automated Planning (1)



- **Deeper discussions** about all of these topics
- **Formal basis** for planning
  - Alternative representations of planning problems
  - Simple and complex state transition systems
- Completely different **principles** for **heuristics**
  - Landmarks
  - Pattern databases
- Alternative **search spaces**
  - Planning by conversion to boolean satisfiability problems
- Extended **expressivity**
  - Planning with time and concurrency
  - Planning with probabilistic actions (Markov Decision Processes)
  - Planning with non-classical goals

- Planning with **domain knowledge**
    - Using what you know: Temporal control rules
    - Breaking down a task into smaller parts: Hierarchical Task Networks



- Alternative **types** of planning
    - Path planning
  - And so on...

